

# **Python AI: How to Build a Neural Network & Make Predictions**

# The Goal of Artificial Intelligence

- Make predictions

$$2 + 2 = 5$$

$$2 + 2 = 4$$

# The Goal of Artificial Intelligence



- Make predictions, that mimic human response
- Human response is not the same as human intelligence

# Predicting the Sum

```
def aisum(a, b):  
    if a == 0: return b  
    if b == 0: return a  
    if a == 1:  
        if b == 1:  
            return 2  
        elif b == 2:  
            return 3  
    elif a == 2:  
        if b == 1:  
            return 3  
        elif b == 2:  
            return 4
```

# Machine Learning

1	0	1
1	1	2
1	2	3
2	0	2
2	2	4

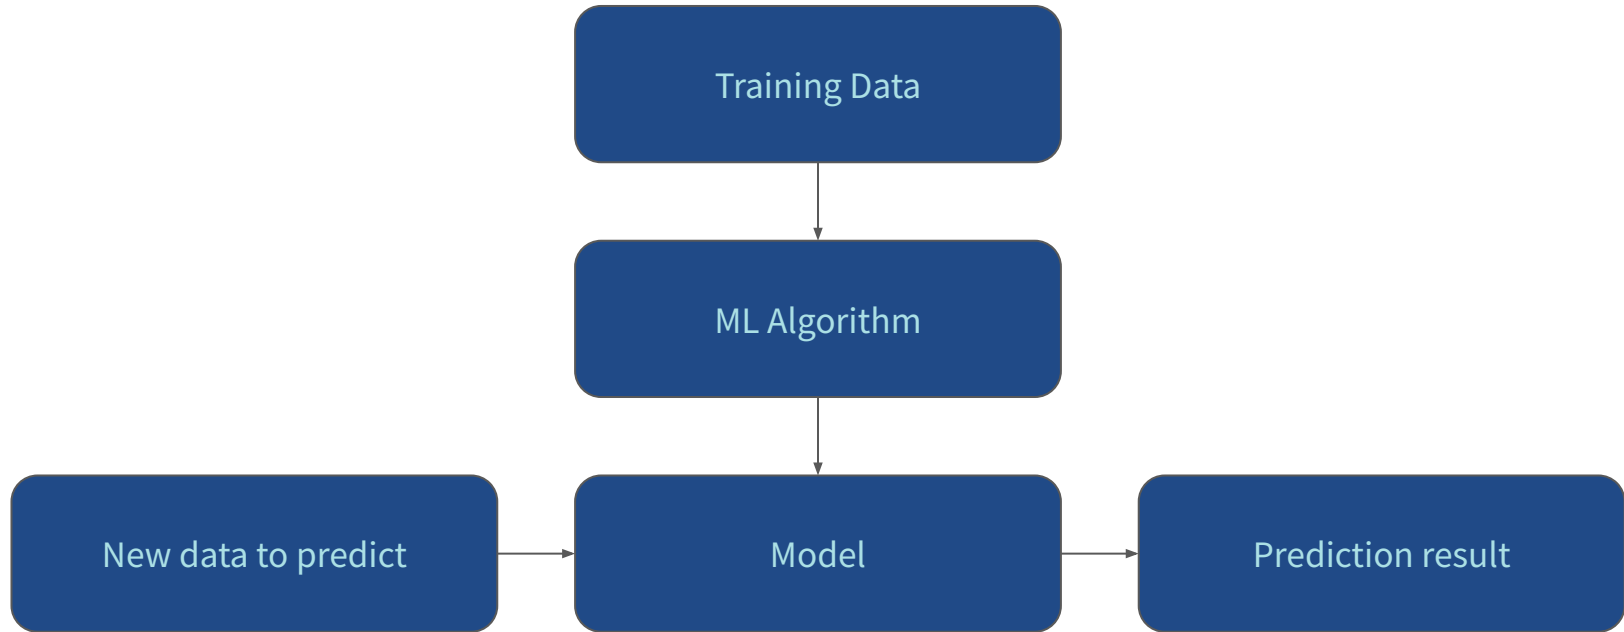


# The Goal of Machine Learning

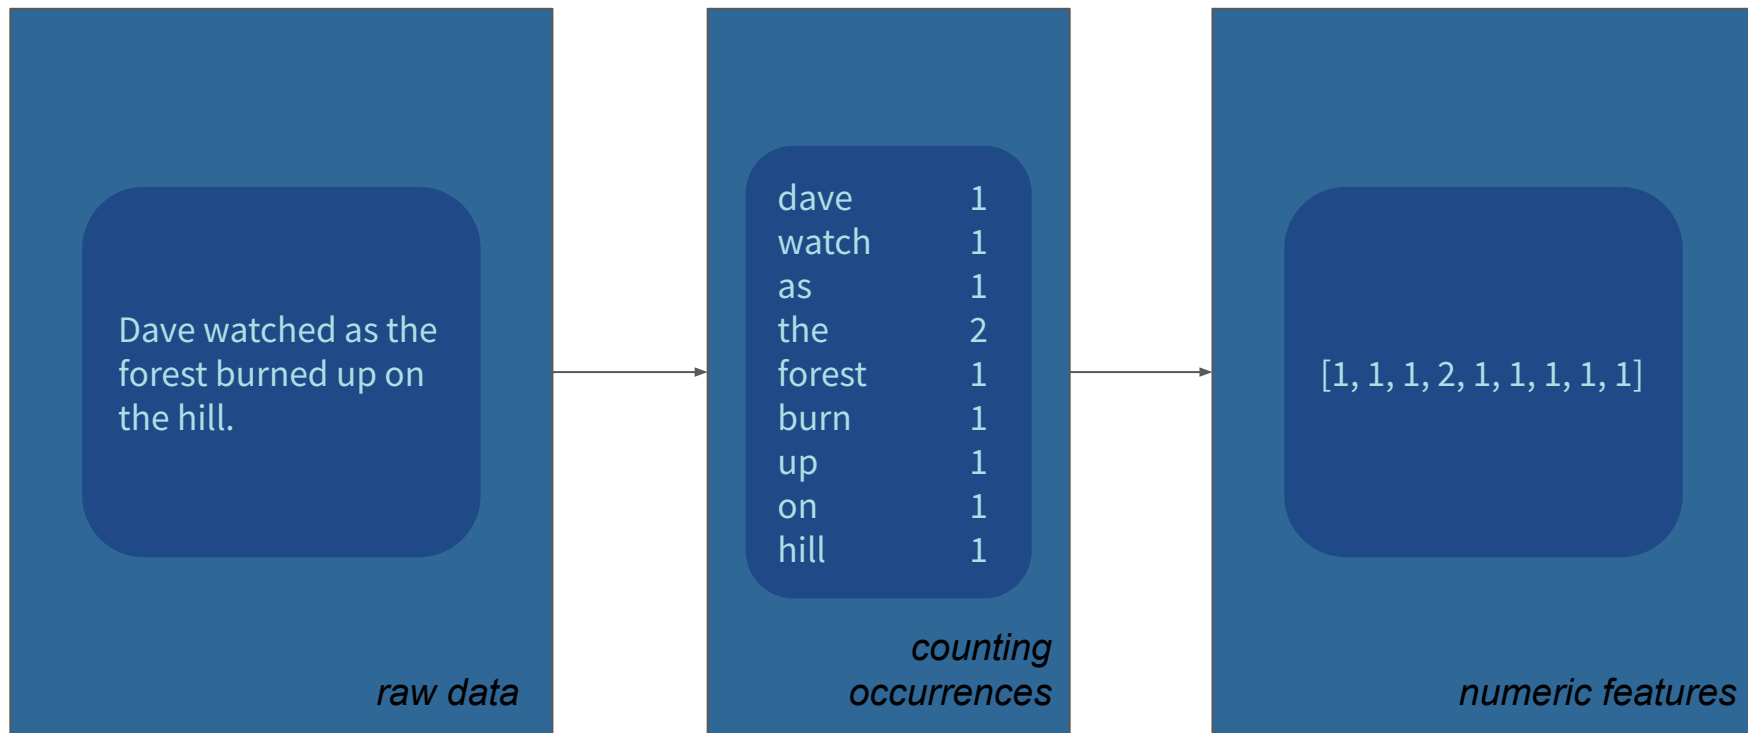
- Make predictions, that mimic human response
- Human response is not the same as human intelligence
- Without being explicitly programmed



# Machine Learning



# Feature Engineering




# The Goal of Machine Learning

- Make predictions, that mimic human response
- Human response is not the same as human intelligence
- Without being explicitly programmed
- Deep learning applies neural networks to machine learning
- Automates feature engineering

# Neural Networks

- Values arranged into layers
- Each layer manipulates the training data
- Output of a layer is the input to the next layer
- The layers are steps in feature engineering


# Neural Networks



# Neural Networks


- Values arranged into layers
- Each layer manipulates the training data
- Output of a layer is the input to the next layer
- The layers are steps in feature engineering
- Vectors
  - numpy ndarray
- Linear regression
  - model the relationship between variables

```
price = (weights_area * area) + (weights_age * age) + bias
```







Input Vector	Target
[1.66, 1.56]	1
[2, 1.5]	0

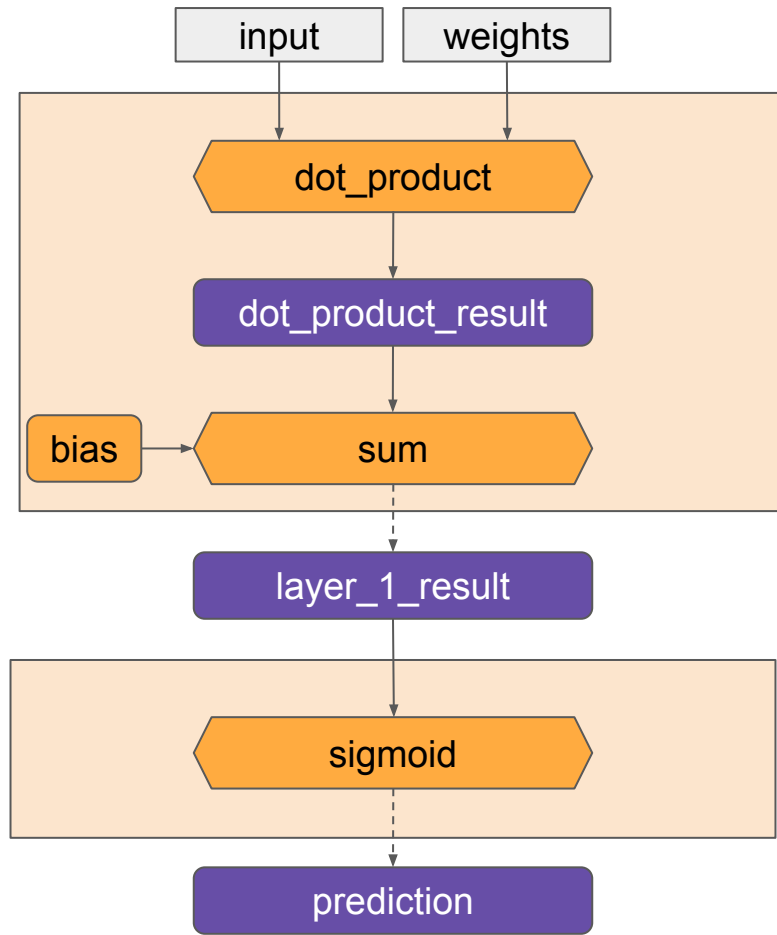


Rectified Linear Unit (ReLU)




Sigmoid





# How Good Is The Prediction?

- Cost or loss function
- Simple mean squared error (MSE)
- $(\text{prediction} - \text{target})^2$



# Computing the Derivative

- Use the derivative to update the values in the network
- Derivative is also called the 'gradient'
- Adjust the weights with the 'gradient descent' process



## Power Rule

$$D(x^n) = n \cdot x^{n-1}$$

$$D(x^2) = 2 \cdot x^{2-1}$$

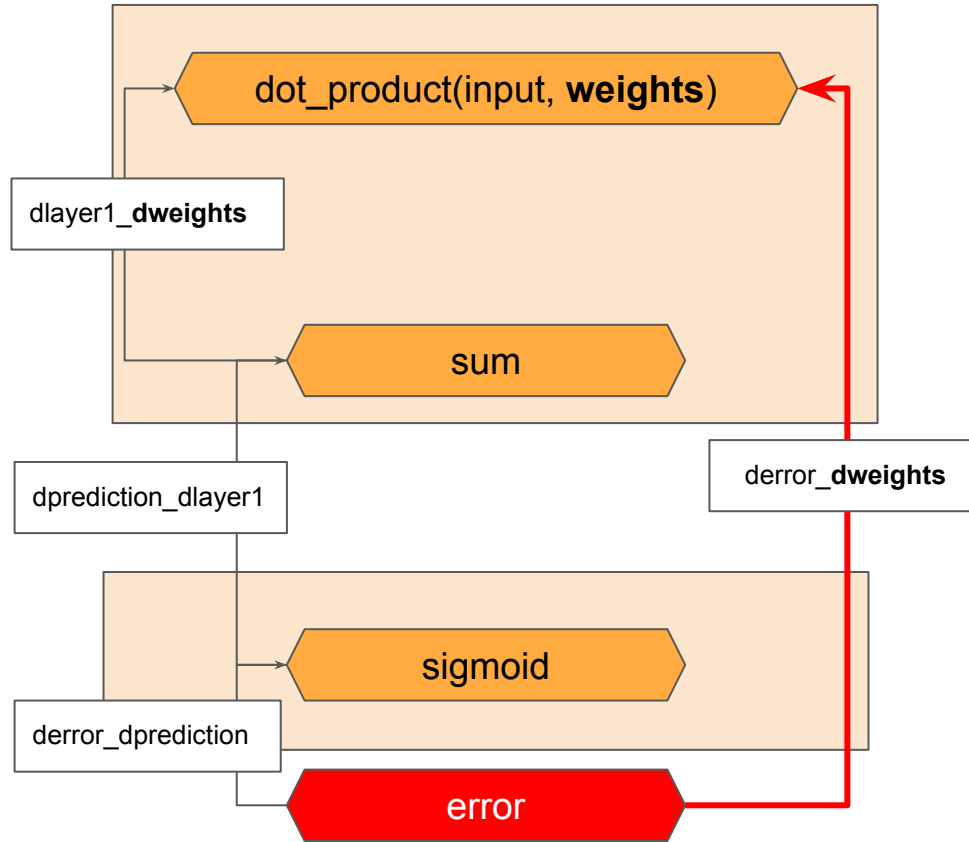
$$D(x^2) = 2 \cdot x^1$$

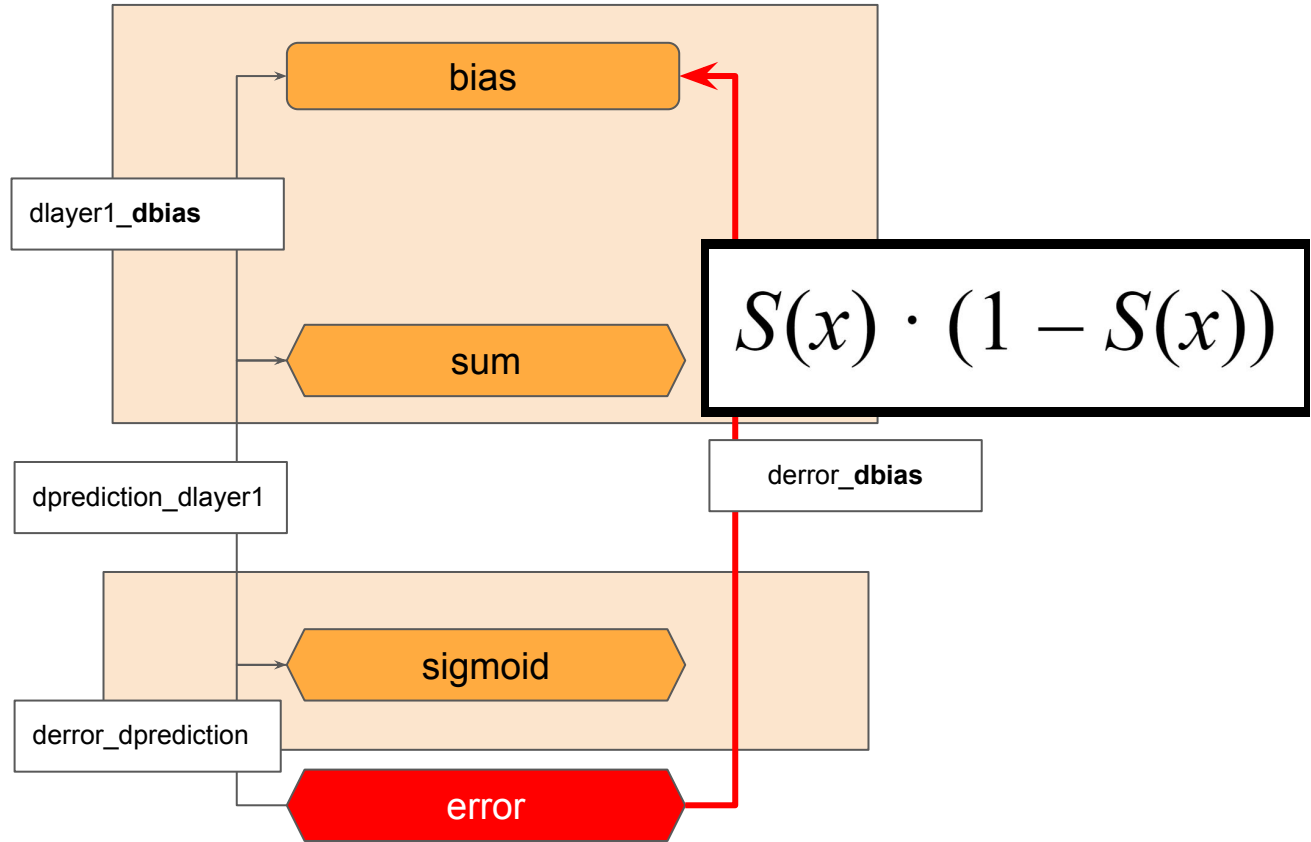
$$D(x^2) = 2x$$

# Function Composition

$$\textit{error} = x^2$$

$$x = \textit{prediction} - \textit{target}$$





```
def sigmoid_deriv(x):  
    return sigmoid(x) * (1 - sigmoid(x))
```

```
derror_prediction = 2 * (prediction - target)  
layer_1 = np.dot(input_vector, weights_1) + bias  
dprediction_dlayer1 = sigmoid_deriv(layer_1)  
dlayer1_dbias = 1
```

```
derror_dbias = (  
    derror_prediction * dprediction_dlayer1 * dlayer1_dbias  
)
```

# Summary

- Foundations of deep learning and neural networks
- Create a neural network in Python
- Deep learning vs. machine learning
- numpy
- Activation functions
- Backpropagation
- Train a model using a neural network
- TensorFlow, Keras and PyTorch

# Resources

- Look Ma, No For-Loops: Array Programming With NumPy
  - <https://realpython.com/numpy-array-programming/>
- Linear Regression in Python
  - <https://realpython.com/linear-regression-in-python/>
- Practical Text Classification With Python and Keras
  - <https://realpython.com/python-keras-text-classification/>
- Pure Python vs NumPy vs TensorFlow Performance Comparison
  - <https://realpython.com/numpy-tensorflow-performance/>
- PyTorch vs TensorFlow for Your Python Deep Learning Project
  - <https://realpython.com/pytorch-vs-tensorflow/>

**THANK YOU!**