

Appendix 4

Implementation of MLP Algorithm and Development and Training of MLP Model

```
[1]: import pandas as pd
import numpy as np
from sklearn.metrics import *
from sklearn.model_selection import train_test_split

import seaborn as sns
import matplotlib.pyplot as plt
import plotly.express as px
import plotly.graph_objects as go
```

1 Loading Datasets

```
[2]: # Loading datas
labels = ['Lagged', 'MA', 'WMA', 'MA-Lagged', 'WMA-Lagged'] # names of each
↳ datasets

def load_datasets():
    """
    Excel files for each dataset are read into a
    dataframe and stored in a dictionary for easy
    access and use
    """
    datasets = dict()
    for lb in labels:
        new_df = pd.read_excel(f"River-Data-{lb}.xlsx")
        new_df.drop(["Unnamed: 0"], axis=1, inplace=True)
        datasets[lb] = new_df

    return datasets

data = load_datasets() # a dataframe for each dataset is stored in a dictionary
↳ called data
```

2 Utility Functions

Standardising and Unstandardising Values

```
[3]: # Utility Functions
## Functions for standardising and unstandardising values
def standardise_columns(df, cols):
    """
    This function works with dataframes to standardise values
    in multiple columns to the range [0.1, 0.9]
    """
    subset_df = df[cols]
    subset_df = 0.8 * ((subset_df - subset_df.min()) / (subset_df.max() -
    ↪subset_df.min())) + 0.1
    return subset_df

def unstandardise_columns(df, cols, max_val, min_val):
    """
    This function works with numpy arrays to destandardise values
    in multiple columns
    """
    subset_df = df[cols]
    subset_df = ((subset_df - subset_df.min()) / 0.8) * (max_val - min_val) +
    ↪min_val
    return subset_df

def standardise_value(x, max_val, min_val):
    """
    This function works with numpy arrays to standardise values
    in multiple arrays to the range [0.1, 0.9]
    """
    return 0.8 * ((x - min_val)) / (max_val - min_val) + 0.1

def unstandardise_value(x, max_val, min_val):
    """
    This function works with numpy arrays to destandardise values
    in multiple arrays
    """
    return ((x - 0.1) / 0.8) * (max_val - min_val) + min_val
```

Plotting

```
[4]: ## Plotting functions
def plot_correlation_matrix(corr_data, title, figsize=(16,6), mask=False):
    """
    Utility function for plotting a correlation heatmap of a given feature set
    """
    if mask:
        mask = np.triu(np.ones_like(corr_data, dtype=bool))
    plt.figure(figsize=figsize, dpi=500)
    heatmap = sns.heatmap(corr_data, vmin=-1, vmax=1, annot=True, mask=mask)
    heatmap.set_title(title)
```

```

plt.show()

def plot_predictions(preds_df, standardised=False):
    """
    Utility function for plotting model predictions against actual value
    """
    preds_col = "Predicted Values"
    vals_col = "Actual Values"
    if standardised:
        preds_col += " (Standardised)"
        vals_col += " (Standardised)"

    line_plt = px.line(preds_df, y=vals_col)
    scatter_plt = px.scatter(preds_df, y=preds_col,
    ↪color_discrete_sequence=["#ff0000"])

    go.Figure(line_plt.data + scatter_plt.data, layout={"title": "Actual vs_
    ↪Predicted Values"}).show()

```

3 ANN Class

```

[5]: # Basic ANN class for MLP models
class BasicAnn:
    def __init__(self, layers, max_st_val, min_st_val, activ_func="sigmoid"):
        self.layers = layers
        self.num_layers = len(layers)
        self.max_val = max_st_val
        self.min_val = min_st_val
        self.activ_func = activ_func

        weight_shapes = [(layers[i-1],layers[i]) for i in range(1, len(layers))]
        self.weights = {
            f"W{i+1}": np.random.standard_normal(s)/s[0]**0.5
            for i, s in enumerate(weight_shapes)
        } # weights are stored as matrices that are implemented as 2D numpy
    ↪arrays

        self.biases = {
            f"B{i+1}": np.random.randn(1,1)/1**0.5
            for i, l in enumerate(layers[1:])
        } # biases are also stored as matrices that are implemented as 2D numpy
    ↪arrays

    def activation(self, x):
        """
        Function to return value with the selected activation
        """

```

```

        if self.activ_func == "sigmoid":
            return 1/(1+np.exp(-x))
        elif self.activ_func == "tanh":
            return (np.exp(x)-np.exp(-x))/(np.exp(x)+np.exp(-x))
        elif self.activ_func == "relu":
            return x * (x > 0)
        elif self.activ_func == "linear":
            return x

def activation_deriv(self, a):
    """
    Function to return value with the derivative of the selected activation
    """
    if self.activ_func == "sigmoid":
        return a * (1 - a)
    elif self.activ_func == "tanh":
        return 1 - a**2
    elif self.activ_func == "relu":
        return 1 * (a > 0)
    elif self.activ_func == "linear":
        return np.ones(a.shape)

def train(self, features, targets, epochs=1000, learning_rate=0.1,
    ↪val_set=None):
    """
    Function will train the model using the standard backpropagation
    ↪algorithm
    and return a dataframe storing various error metrics for the model on
    ↪the
    training set and, possibly, a validation set if that is given
    """
    results = pd.DataFrame()
    real_targets = unstandardise_value(targets, self.max_val, self.min_val)
    num_targets = len(targets)

    for _ in range(epochs):
        # Forward pass
        activations = self.forward_pass(features)

        # Error calculation
        output_layer = activations[f"A{self.num_layers - 1}"]
        real_preds = unstandardise_value(output_layer, self.max_val, self.
    ↪min_val)
        error_data = { # storing error metrics for both standardised and
    ↪unstandardised data
            "mse": mean_squared_error(real_targets, real_preds),

```

```

        "rmse": mean_squared_error(real_targets, real_preds,
→squared=False),
        "mae": mean_absolute_error(real_targets, real_preds),
        "r_sqr": r2_score(real_targets, real_preds),
        "st_mse": mean_squared_error(targets, output_layer),
        "st_rmse": mean_squared_error(targets, output_layer,
→squared=False),
        "st_mae": mean_absolute_error(targets, output_layer),
        "st_r_sqr": r2_score(targets, output_layer)
    }

    if val_set:
        # if there is a validation set the prediction error of the model
        # on the validation set will be stored
        r, err = self.predict(val_set[0].to_numpy(), val_set[1].
→to_numpy())
        error_data.update({f"val_{col}": err[col][0] for col in err.
→columns})

    results = results.append(error_data, ignore_index=True)

    # Backward pass (backpropagation algorithm)
    deltas = self.compute_deltas(activations, targets, output_layer)
    self.update_weights(deltas, activations, features, num_targets,
→learning_rate)

    return results

def predict(self, test_inputs, st_actual_outputs, actual_outputs=None):
    """
    Runs a forward pass of the network with the newly configured weights
    and biases and returns a dataframe comparing the predicted values
    to actual values as well as a dataframe with various error metrics
    """
    # Forward pass
    activations = self.forward_pass(test_inputs)
    st_preds = activations[f"A{self.num_layers - 1}"]

    # Comparing predicted values with actual values
    if actual_outputs is None:
        actual_outputs = unstandardise_value(st_actual_outputs, self.
→max_val, self.min_val)

    preds = unstandardise_value(st_preds, self.max_val, self.min_val)

    results = pd.DataFrame(

```

```

        data={
            "Actual Values": actual_outputs.flatten(),
            "Predicted Values": preds.flatten(),
            "Actual Values (Standardised)": st_actual_outputs.flatten(),
            "Predicted Values (Standardised)": st_preds.flatten(),
        }
    )

    # Error calculation
    results["Absolute Error"] = abs(results["Actual Values"] -
↪results["Predicted Values"])
    st_absolute_err = abs(results["Actual Values (Standardised)"] -
↪results["Predicted Values (Standardised)"])
    results["Absolute Error (Standardised Values)"] = st_absolute_err

    error_metrics = pd.DataFrame(data={
        "mse": [mean_squared_error(actual_outputs, preds)],
        "rmse": [mean_squared_error(actual_outputs, preds, squared=False)],
        "mae": [mean_absolute_error(actual_outputs, preds)],
        "r_sqr": [r2_score(actual_outputs, preds)],
        "st_mse": [mean_squared_error(st_actual_outputs, st_preds)],
        "st_rmse": [mean_squared_error(st_actual_outputs, st_preds,
↪squared=False)],
        "st_mae": [mean_absolute_error(st_actual_outputs, st_preds)],
        "st_r_sqr": [r2_score(st_actual_outputs, st_preds)]
    })

    return results, error_metrics

def forward_pass(self, features):
    """
    Runs a forward pass of neural network through repeated
    multiplication of weights and bias matrices. Returns
    list of each activation layer including the output layer.
    """
    activation = self.activation(np.dot(features, self.weights["W1"]) +
↪self.biases["B1"].T)
    activations = {"A1": activation}
    for i in range(2, self.num_layers):
        activation = self.activation(np.dot(activation, self.
↪weights[f"W{i}"]) + self.biases[f"B{i}"].T)
        activations[f"A{i}"] = activation

    return activations

def compute_deltas(self, activations, targets, output_layer):
    """

```

```

Computes errors between layers for backproagation.
Returns a dictionary of lists which contain the errors
for each node in each layer.
"""
output_err = targets - output_layer
output_delta = output_err * self.activation_deriv(output_layer)
deltas = {"dw1": output_delta}

for i in range(self.num_layers - 1, 1, -1):
    dw = deltas[f"dw{self.num_layers - i}"]
    act = activations[f"A{i-1}"]
    w = self.weights[f"W{i}"]
    deltas[f"dw{self.num_layers - i + 1}"] = np.dot(dw, w.T) * self.
↪activation_deriv(act)

    return deltas

def update_weights(self, deltas, activations, features, num_targets,
↪l_rate):
    """
    Updates weights and biases according to given errors, activations
    and the chosen learning rate
    """
    delta = deltas[f"dw{self.num_layers - 1}"]
    self.weights["W1"] += l_rate * (np.dot(features.T, delta)) / num_targets
    self.biases["B1"] += l_rate * (np.dot(delta.T, np.ones((num_targets,
↪1)))) / num_targets

    for i in range(2, self.num_layers):
        act = activations[f"A{i-1}"]
        dw = deltas[f"dw{self.num_layers - i}"]
        self.weights[f"W{i}"] += l_rate * (np.dot(act.T, dw)) / num_targets
        self.biases[f"B{i}"] += l_rate * np.dot(dw.T, np.ones((num_targets,
↪1)))) / num_targets

```

Build, Train and Test ANN Model

```

[6]: def build_train_test(feature_set, feature_cols, target_cols, layers=("auto",
↪1), activ_func="linear", epochs=1000, l_rate=0.1):
    """
    Function to build, train and test MLP models
    """
    # Splitting and standardising datasets to create standardised and
↪unstandardised
    # training, validation and testing sets.
    train_val_set, test_set = train_test_split(feature_set, test_size=0.2)
    st_train_val_set = standardise_columns(train_val_set, train_val_set.columns)

```

```

st_test_set = standardise_columns(test_set, test_set.columns)

# Preparing features and targets for training and testing
features = st_train_val_set[feature_cols]
targets = st_train_val_set[target_cols]

X_train, X_val, y_train, y_val = train_test_split(features, targets,
→test_size=0.25)
X_test, y_test = st_test_set[feature_cols], st_test_set[target_cols]

# Getting standardisation values for targets
min_val = train_val_set[target_cols].min()[0]
max_val = train_val_set[target_cols].max()[0]

# Building model
if layers[0] == "auto":
    # if the size of the input layer is not specified
    # then it will be set to the number of predictors
    layers = (len(feature_cols),) + layers[1:]

ann = BasicAnn(layers, max_val, min_val, activ_func)

# Training model
training_results = ann.train(
    X_train.to_numpy(),
    y_train.to_numpy(),
    val_set=(X_val, y_val), # training with a validation set
    epochs=epochs,
    learning_rate=l_rate
)

# Predicting model
prediction_results = ann.predict(
    X_test.to_numpy(),
    y_test.to_numpy(),
    actual_outputs=test_set[target_cols].to_numpy()
)

predictions, error_metrics = prediction_results[0], prediction_results[1]

return {
    "training_results": training_results,
    "final_test_results": predictions,
    "error_metrics": error_metrics,
    "model": ann
}

```


4 Selecting Features/Predictors

Building Feature Sets

```
[7]: # Function for building custom feature and target sets
def build_feature_set(*datasets):
    assert len(datasets) > 0, "No data sets entered"
    datasets = list(datasets)
    min_rows = min(d.shape[0] for d in datasets)

    for i, ds in enumerate(datasets):
        datasets[i] = ds.truncate(before=ds.shape[0]-min_rows).reset_index()
        datasets[i].drop(["index"], axis=1, inplace=True)

    merged_df = datasets[0].iloc[:, :2]
    for ds in datasets:
        merged_df = pd.concat([merged_df, ds.iloc[:, 2:]], axis=1)

    merged_cols = list(merged_df.columns)
    selected_cols = []

    for i in range(0, len(merged_cols), 2):
        format_str = f"{i+1}) {merged_cols[i]}"
        if i != len(merged_cols) - 1:
            second_part = f"{i+2}) {merged_cols[i+1]}"
            num_spaces = 50 - len(format_str)
            format_str += num_spaces*" " + second_part
        print(format_str)

    selected_indices = input("\nSelect columns: ")
    for index in selected_indices.split(","):
        if "-" in index:
            first_i, second_i = index.split("-")
            selected_cols += merged_cols[int(first_i) - 1: int(second_i)]
        else:
            selected_cols.append(merged_cols[int(index) - 1])

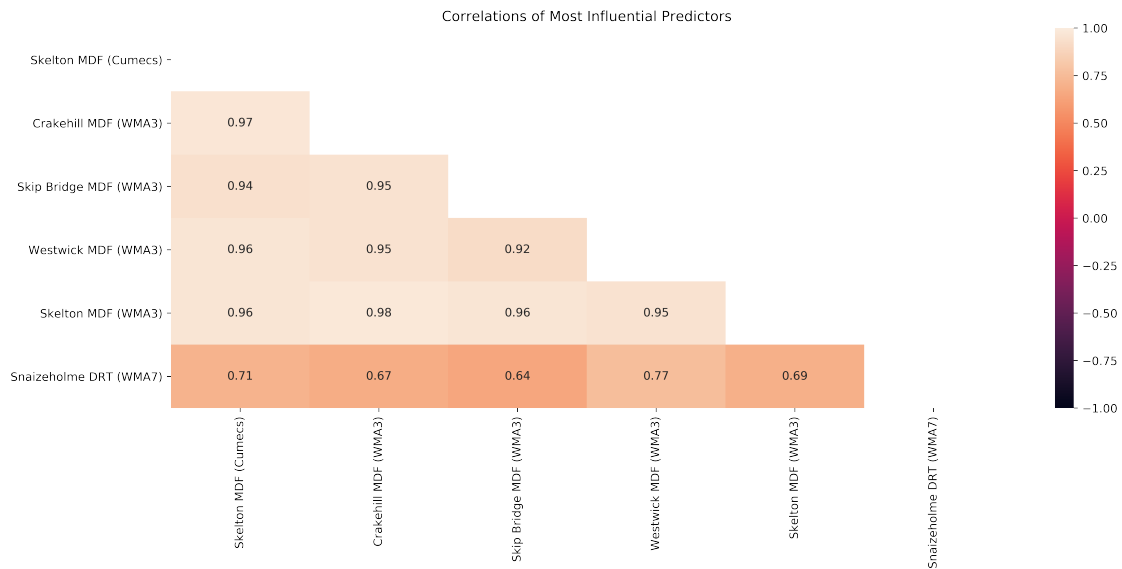
    return merged_df[selected_cols]
```

```
[8]: # 2,3-6,42
fs = build_feature_set(data['WMA'])
plot_correlation_matrix(fs.corr(), "Correlations of Most Influential_
↳Predictors", mask=True)
```

- | | |
|-------------------------|---------------------------|
| 1) Date | 2) Skelton MDF (Cumeecs) |
| 3) Crakehill MDF (WMA3) | 4) Skip Bridge MDF (WMA3) |
| 5) Westwick MDF (WMA3) | 6) Skelton MDF (WMA3) |
| 7) Crakehill MDF (WMA4) | 8) Skip Bridge MDF (WMA4) |
| 9) Westwick MDF (WMA4) | 10) Skelton MDF (WMA4) |

11) Crakehill MDF (WMA5)	12) Skip Bridge MDF (WMA5)
13) Westwick MDF (WMA5)	14) Skelton MDF (WMA5)
15) Crakehill MDF (WMA6)	16) Skip Bridge MDF (WMA6)
17) Westwick MDF (WMA6)	18) Skelton MDF (WMA6)
19) Crakehill MDF (WMA7)	20) Skip Bridge MDF (WMA7)
21) Westwick MDF (WMA7)	22) Skelton MDF (WMA7)
23) Arkengarthdale DRT (WMA3)	24) East Cowton DRT (WMA3)
25) Malham Tarn DRT (WMA3)	26) Snaizeholme DRT (WMA3)
27) Arkengarthdale DRT (WMA4)	28) East Cowton DRT (WMA4)
29) Malham Tarn DRT (WMA4)	30) Snaizeholme DRT (WMA4)
31) Arkengarthdale DRT (WMA5)	32) East Cowton DRT (WMA5)
33) Malham Tarn DRT (WMA5)	34) Snaizeholme DRT (WMA5)
35) Arkengarthdale DRT (WMA6)	36) East Cowton DRT (WMA6)
37) Malham Tarn DRT (WMA6)	38) Snaizeholme DRT (WMA6)
39) Arkengarthdale DRT (WMA7)	40) East Cowton DRT (WMA7)
41) Malham Tarn DRT (WMA7)	42) Snaizeholme DRT (WMA7)

Select columns: 2,3-6,42



5 Training and Network Selection

Epochs

```
[9]: target_cols = [fs.columns[0]]
    feature_cols = list(fs.columns[1:])

    epoch_tests = dict()
    for i in range(1, 11):
        epoch_tests[f"Test-{i}"] = build_train_test(
```

```

        fs,
        feature_cols,
        target_cols,
        layers=("auto", 1),
        activ_func="linear",
        epochs=i*500
    )

```

- rmse -> root mean square error
- mae -> mean absolute error
- mse -> mean squared error
- r_sqr -> R-Squared (Coefficient of Determination)
- val_* -> error metric on validation set
- st_* -> error metric on standardised values

```

[10]: # Testing number of epochs for training; between 800 and 1200 seems to be ideal
for i in range(1, 11):
    print(f"Model trained with {i*500} epochs", end=f"\n{'-'*100}\n")

    print("Final Training results")
    print(epoch_tests[f"Test-{i}"]["training_results"].iloc[-1, :4],
    ↪end=f"\n{'-'*100}\n")

    print("Final Validation results")
    print(epoch_tests[f"Test-{i}"]["training_results"].iloc[-1, 8:12],
    ↪end=f"\n{'-'*100}\n")

    print("Test Set Results")
    print(epoch_tests[f"Test-{i}"]["error_metrics"].iloc[0],
    ↪end=f"\n{'-'*100}\n\n\n")

    ax = epoch_tests[f"Test-{i}"]["training_results"].plot(
        y=["rmse", "val_rmse"], title=f"Model Trained with {i*500} Epochs",
    )
    ax.set_xlabel("Epochs")
    ax.set_ylabel("Root Mean Squared Error (RMSE)")

```

Model trained with 500 epochs

```

-----
Final Training results
mae      11.115425
mse      298.617903
r_sqr     0.907014
rmse     17.280564

```

Name: 499, dtype: float64

Final Validation results

val_mae 11.824059

val_mse 355.111826

val_r_sqr 0.899109

val_rmse 18.844411

Name: 499, dtype: float64

Test Set Results

mse 540.297189

rmse 23.244294

mae 18.232164

r_sqr 0.766672

st_mse 0.001970

st_rmse 0.044380

st_mae 0.025539

st_r_sqr 0.888120

Name: 0, dtype: float64

=====

Model trained with 1000 epochs

Final Training results

mae 8.688907

mse 212.143011

r_sqr 0.937055

rmse 14.565130

Name: 999, dtype: float64

Final Validation results

val_mae 8.612176

val_mse 224.111494

val_r_sqr 0.928040

val_rmse 14.970354

Name: 999, dtype: float64

Test Set Results

mse 263.936692

rmse 16.246129

```
mae          9.433544
r_sqr        0.880873
st_mse       0.000677
st_rmse      0.026011
st_mae       0.016009
st_r_sqr     0.938698
Name: 0, dtype: float64
```

```
=====
=====
```

Model trained with 1500 epochs

```
-----
-----
```

Final Training results

```
mae          7.725300
mse         180.929892
r_sqr        0.939766
rmse        13.451018
Name: 1499, dtype: float64
```

```
-----
-----
```

Final Validation results

```
val_mae       8.400230
val_mse       230.727938
val_r_sqr     0.930766
val_rmse      15.189731
Name: 1499, dtype: float64
```

```
-----
-----
```

Test Set Results

```
mse          763.629961
rmse         27.633855
mae          17.428569
r_sqr        0.756392
st_mse       0.001362
st_rmse      0.036908
st_mae       0.019721
st_r_sqr     0.933314
Name: 0, dtype: float64
```

```
=====
=====
```

Model trained with 2000 epochs

```
-----
```

```
-----  
Final Training results  
mae          9.859058  
mse          323.622713  
r_sqr        0.903529  
rmse         17.989517  
Name: 1999, dtype: float64  
-----
```

```
-----  
Final Validation results  
val_mae       9.298527  
val_mse       262.621712  
val_r_sqr     0.897150  
val_rmse      16.205607  
Name: 1999, dtype: float64  
-----
```

```
-----  
Test Set Results  
mse          1534.300343  
rmse         39.170146  
mae          23.837348  
r_sqr        0.465289  
st_mse       0.001491  
st_rmse      0.038611  
st_mae       0.023668  
st_r_sqr     0.919180  
Name: 0, dtype: float64  
=====
```

```
=====
```

Model trained with 2500 epochs

```
-----
```

```
-----  
Final Training results  
mae          10.612285  
mse          311.711768  
r_sqr        0.905703  
rmse         17.655361  
Name: 2499, dtype: float64  
-----
```

```
-----  
Final Validation results  
val_mae       9.807566  
val_mse       261.335346  
val_r_sqr     0.908631  
val_rmse      16.165870
```

Name: 2499, dtype: float64

Test Set Results

mse	2619.121607
rmse	51.177354
mae	32.500738
r_sqr	0.035265
st_mse	0.003658
st_rmse	0.060480
st_mae	0.037223
st_r_sqr	0.766383

Name: 0, dtype: float64

=====

Model trained with 3000 epochs

Final Training results

mae	6.129363
mse	124.789594
r_sqr	0.960619
rmse	11.170926

Name: 2999, dtype: float64

Final Validation results

val_mae	6.326957
val_mse	124.882068
val_r_sqr	0.952840
val_rmse	11.175065

Name: 2999, dtype: float64

Test Set Results

mse	648.318969
rmse	25.462109
mae	14.113753
r_sqr	0.802339
st_mse	0.001672
st_rmse	0.040895
st_mae	0.024533
st_r_sqr	0.926110

Name: 0, dtype: float64

=====

Model trained with 3500 epochs

Final Training results
mae 8.736078
mse 250.379080
r_sqr 0.920063
rmse 15.823371
Name: 3499, dtype: float64

Final Validation results
val_mae 9.033819
val_mse 256.454614
val_r_sqr 0.921544
val_rmse 16.014200
Name: 3499, dtype: float64

Test Set Results
mse 313.532026
rmse 17.706836
mae 10.328859
r_sqr 0.888670
st_mse 0.004751
st_rmse 0.068925
st_mae 0.038536
st_r_sqr 0.740614
Name: 0, dtype: float64

=====

Model trained with 4000 epochs

Final Training results
mae 6.670029
mse 157.370618
r_sqr 0.949555
rmse 12.544745
Name: 3999, dtype: float64


```
-----  
Final Validation results  
val_mae      6.234223  
val_mse     178.456109  
val_r_sqr    0.935996  
val_rmse    13.358747  
Name: 3999, dtype: float64  
-----
```

```
-----  
Test Set Results  
mse          212.698633  
rmse         14.584191  
mae           7.794754  
r_sqr         0.936354  
st_mse        0.000564  
st_rmse       0.023739  
st_mae        0.013208  
st_r_sqr      0.948000  
Name: 0, dtype: float64  
=====
```

```
=====
```

Model trained with 4500 epochs

```
-----  
Final Training results  
mae          6.500001  
mse          155.093720  
r_sqr         0.946919  
rmse         12.453663  
Name: 4499, dtype: float64  
-----
```

```
-----  
Final Validation results  
val_mae      6.386406  
val_mse     131.872036  
val_r_sqr    0.959946  
val_rmse    11.483555  
Name: 4499, dtype: float64  
-----
```

```
-----  
Test Set Results  
mse          932.030656  
rmse         30.529177  
mae          17.138363  
r_sqr         0.725804
```

```
st_mse      0.001296
st_rmse     0.036001
st_mae      0.020815
st_r_sqr    0.940711
Name: 0, dtype: float64
```

```
=====
=====
```

Model trained with 5000 epochs

Final Training results

```
mae      8.351984
mse     199.404995
r_sqr    0.934940
rmse     14.121083
Name: 4999, dtype: float64
```

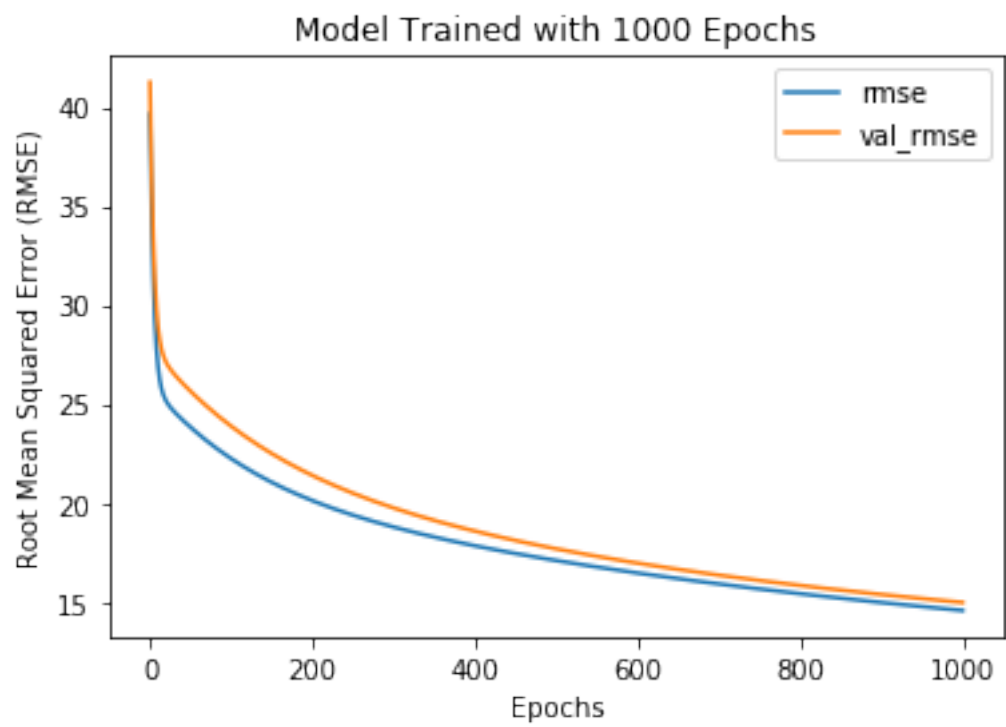
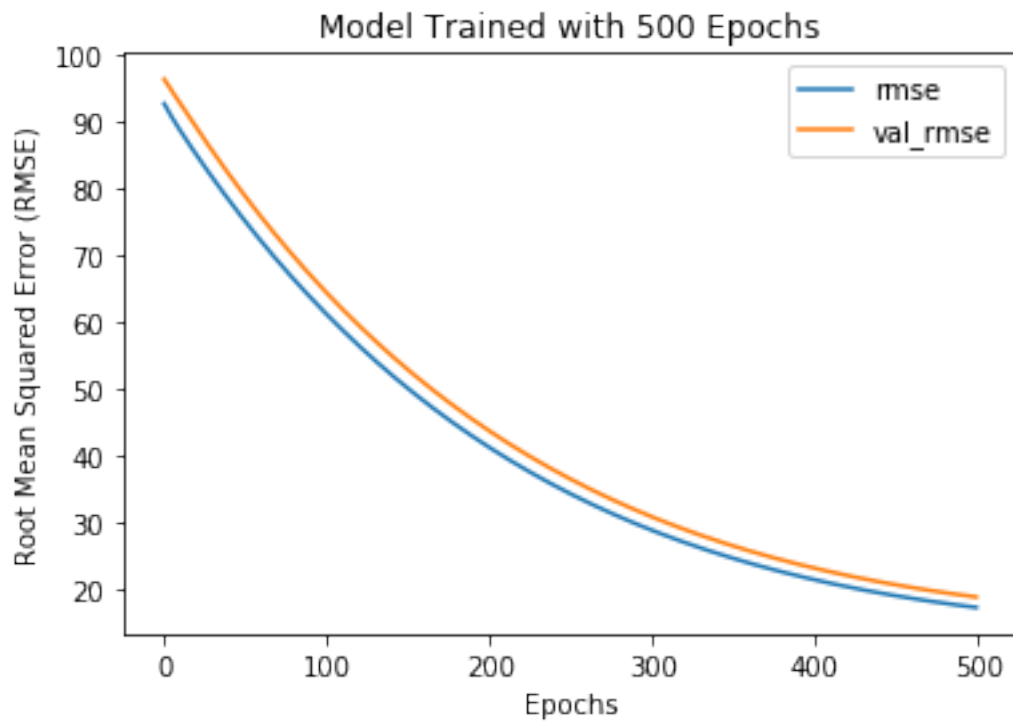
Final Validation results

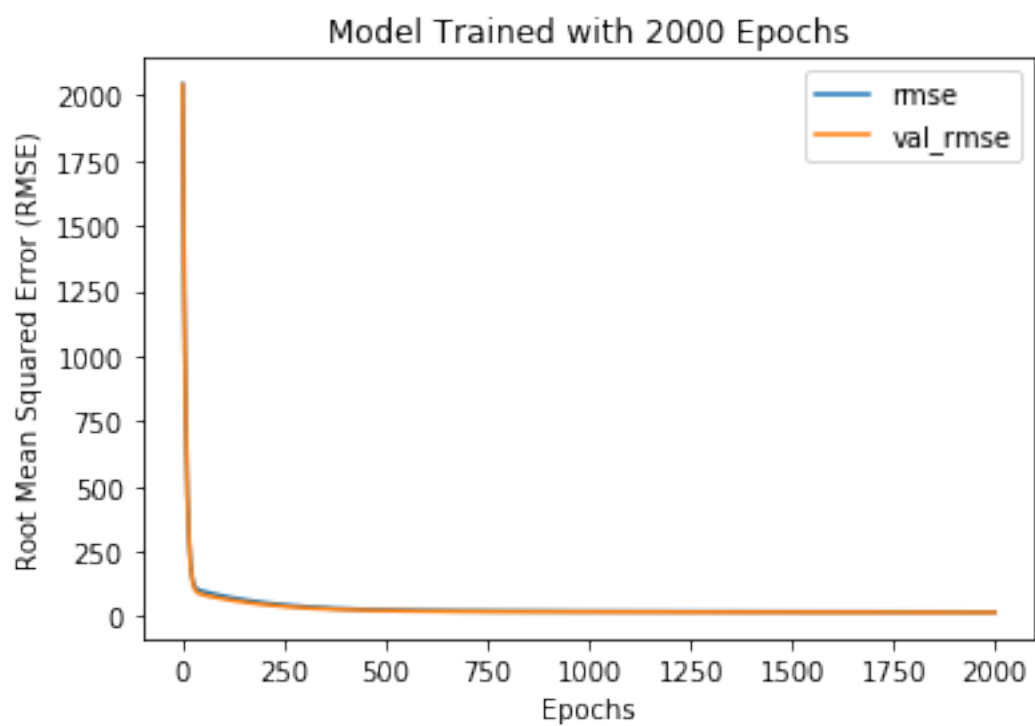
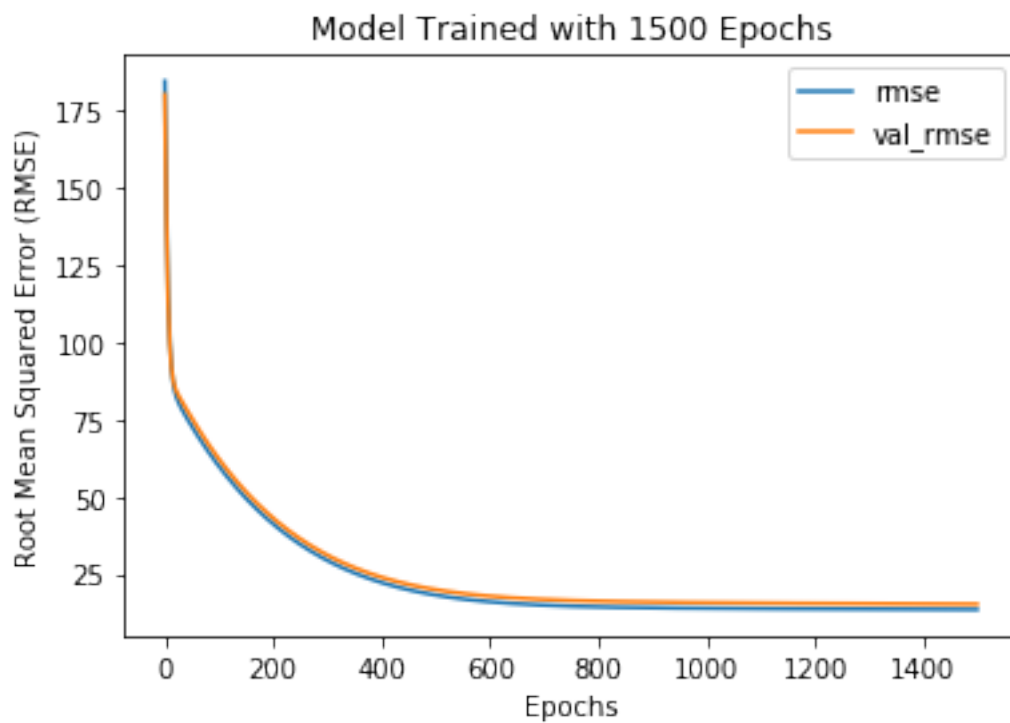
```
val_mae    8.269012
val_mse   176.502604
val_r_sqr  0.949772
val_rmse   13.285428
Name: 4999, dtype: float64
```

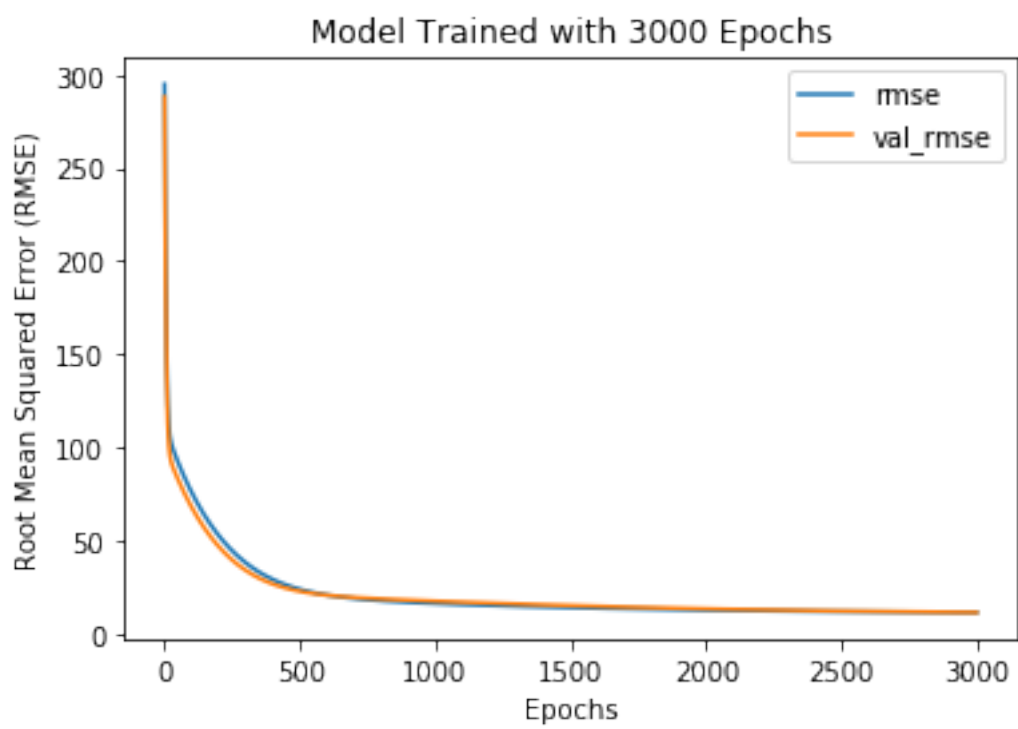
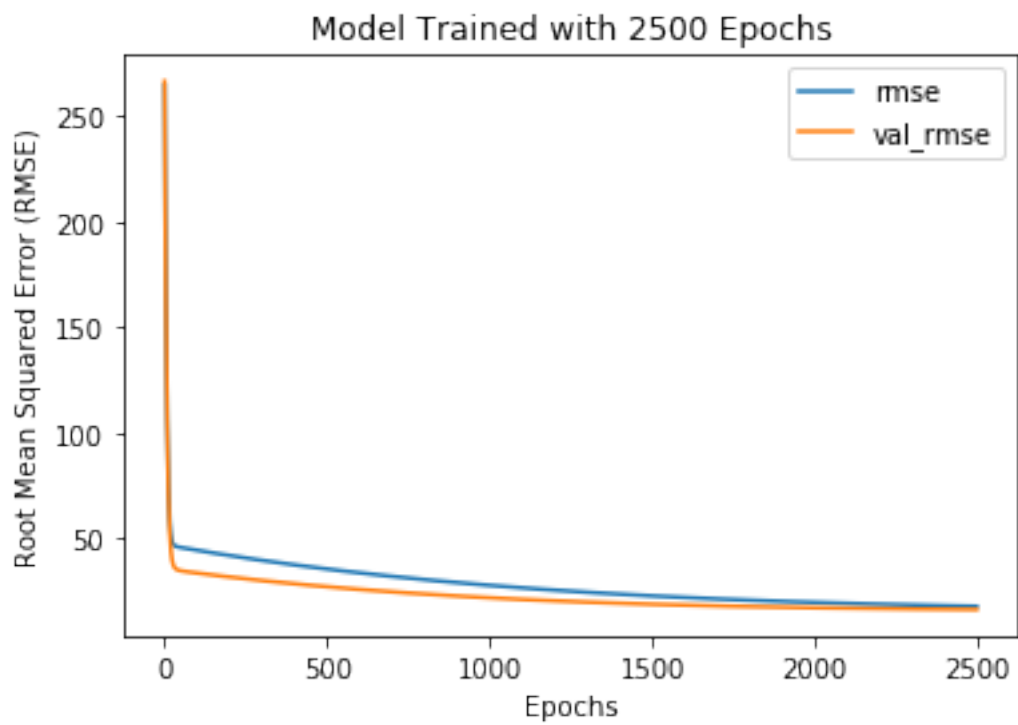
Test Set Results

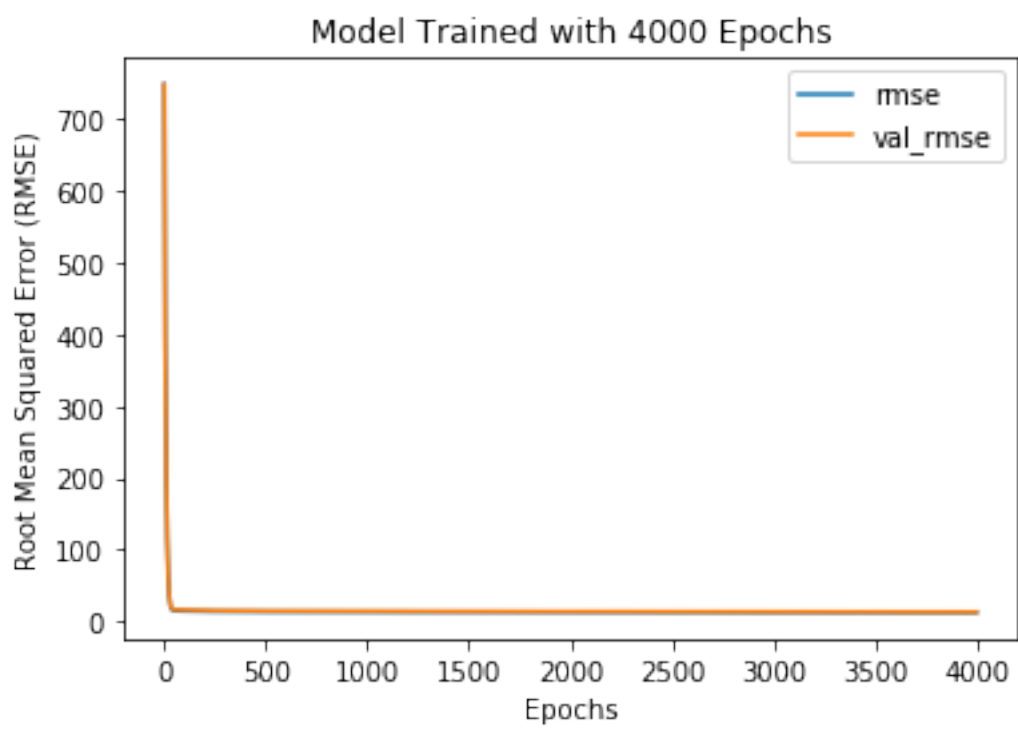
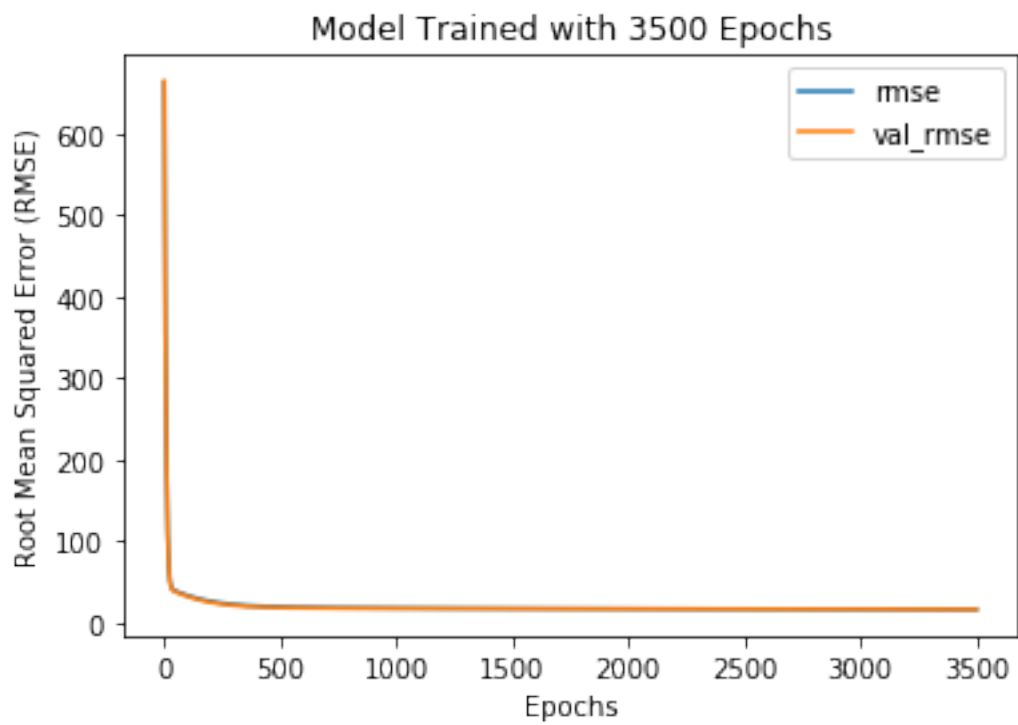
```
mse      830.270766
rmse     28.814419
mae      15.929263
r_sqr    0.702218
st_mse    0.005381
st_rmse   0.073353
st_mae    0.039558
st_r_sqr  0.794596
Name: 0, dtype: float64
```

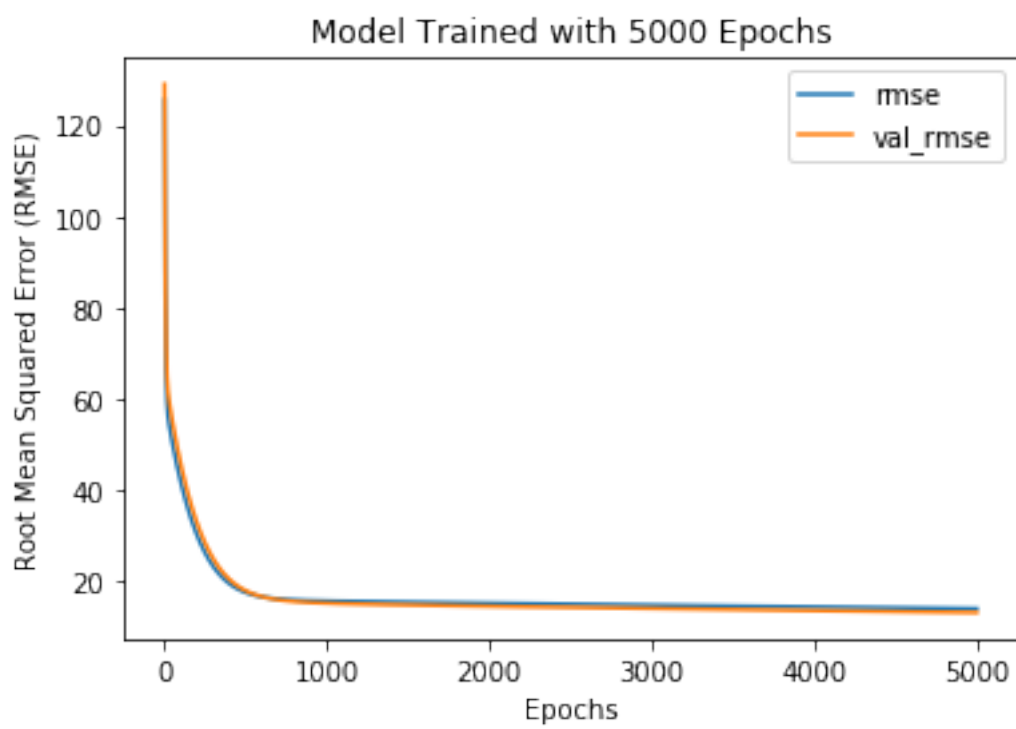
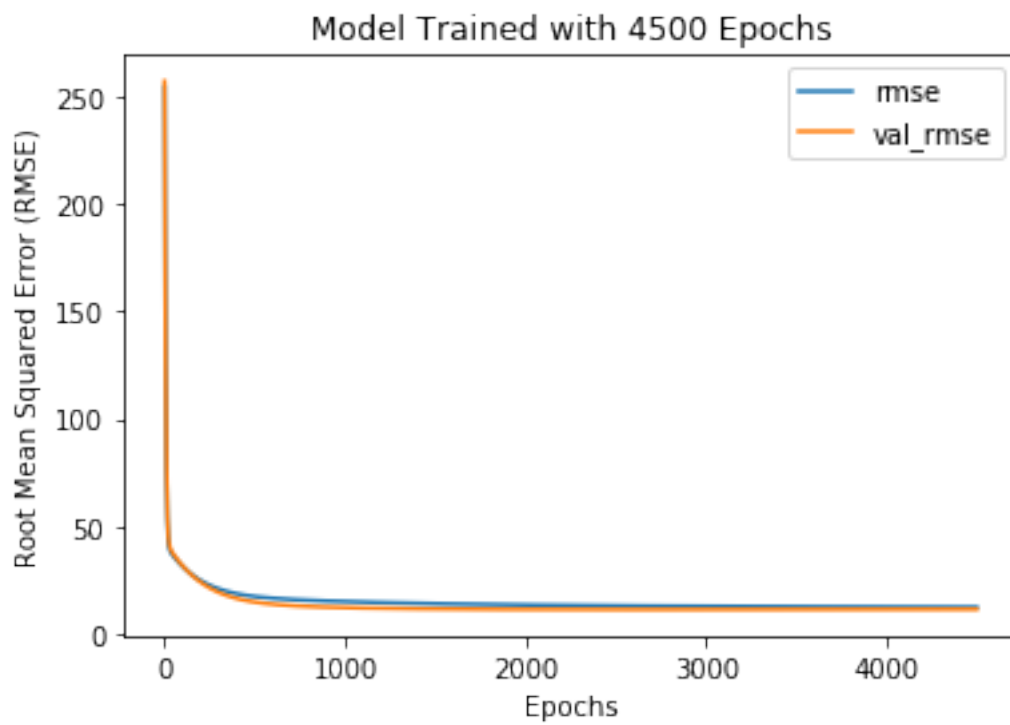
```
=====
=====
```











- Models seem to overfit after about 800 to 1200 epochs
- Suitable number for training is around 1000

Learning Rate

```
[11]: l_rate_tests = dict()
      for i in range(1, 11):
          l_rate_tests[f"Test-{i}"] = build_train_test(
              fs,
              feature_cols,
              target_cols,
              layers=("auto", 1),
              activ_func="linear",
              epochs=1000,
              l_rate=0.1*i
          )

[12]: for i in range(1,11):
      print(f"Model Trained with Learning Rate of {0.1*i}", end=f"\n{'-'*100}\n")

      print("Final Training results")
      print(l_rate_tests[f"Test-{i}"]["training_results"].iloc[-1, :4],
            ↪end=f"\n{'-'*100}\n")

      print("Final Validation results")
      print(l_rate_tests[f"Test-{i}"]["training_results"].iloc[-1, 8:12],
            ↪end=f"\n{'-'*100}\n")

      print("Test Set Results")
      print(l_rate_tests[f"Test-{i}"]["error_metrics"].iloc[0],
            ↪end=f"\n{'-'*100}\n\n\n")

      ax = l_rate_tests[f"Test-{i}"]["training_results"].plot(
          y=["rmse", "val_rmse"], title=f"Model Trained with Learning Rate of
            ↪{i*0.1}",
      )
      ax.set_xlabel("Epochs")
      ax.set_ylabel("Root Mean Squared Error (RMSE)")
```

Model Trained with Learning Rate of 0.1

```
-----
Final Training results
mae      6.728895
mse     144.544700
r_sqr     0.954585
rmse     12.022674
Name: 999, dtype: float64
-----
```



```
-----  
Final Validation results  
val_mae      6.827042  
val_mse      154.903146  
val_r_sqr     0.951899  
val_rmse     12.446009  
Name: 999, dtype: float64  
-----
```

```
-----  
Test Set Results  
mse          166.056780  
rmse         12.886302  
mae          6.526054  
r_sqr        0.938928  
st_mse       0.001043  
st_rmse      0.032299  
st_mae       0.015748  
st_r_sqr     0.923022  
Name: 0, dtype: float64  
=====
```

Model Trained with Learning Rate of 0.2

```
-----  
Final Training results  
mae          9.262914  
mse          288.251083  
r_sqr        0.913986  
rmse         16.977959  
Name: 999, dtype: float64  
-----
```

```
-----  
Final Validation results  
val_mae      8.271171  
val_mse      240.247581  
val_r_sqr     0.909595  
val_rmse     15.499922  
Name: 999, dtype: float64  
-----
```

```
-----  
Test Set Results  
mse          361.894289  
rmse         19.023519  
mae          9.305752  
r_sqr        0.867981
```

```
st_mse      0.004561
st_rmse     0.067532
st_mae      0.035913
st_r_sqr    0.800938
Name: 0, dtype: float64
```

```
=====
=====
```

Model Trained with Learning Rate of 0.30000000000000004

```
-----
-----
```

Final Training results

```
mae      8.066564
mse     204.581414
r_sqr    0.930061
rmse    14.303196
Name: 999, dtype: float64
```

```
-----
-----
```

Final Validation results

```
val_mae    8.145104
val_mse   229.466800
val_r_sqr   0.916561
val_rmse   15.148162
Name: 999, dtype: float64
```

```
-----
-----
```

Test Set Results

```
mse      236.679537
rmse     15.384393
mae       9.438774
r_sqr    0.939067
st_mse    0.001687
st_rmse   0.041069
st_mae    0.024278
st_r_sqr  0.912746
Name: 0, dtype: float64
```

```
=====
=====
```

Model Trained with Learning Rate of 0.4

```
-----
-----
```

Final Training results

```
mae      8.032468
mse      223.892823
r_sqr    0.926797
rmse     14.963049
Name: 999, dtype: float64
```

Final Validation results

```
val_mae      8.067513
val_mse     192.114710
val_r_sqr    0.934863
val_rmse     13.860545
Name: 999, dtype: float64
```

Test Set Results

```
mse      359.787579
rmse     18.968067
mae      11.404283
r_sqr    0.892779
st_mse    0.001621
st_rmse   0.040257
st_mae    0.021636
st_r_sqr  0.916093
Name: 0, dtype: float64
```

```
=====
=====
```

Model Trained with Learning Rate of 0.5

Final Training results

```
mae      6.481272
mse     137.914126
r_sqr    0.955400
rmse     11.743685
Name: 999, dtype: float64
```

Final Validation results

```
val_mae      6.547086
val_mse     172.664087
val_r_sqr    0.950973
val_rmse     13.140171
Name: 999, dtype: float64
```

```

-----
Test Set Results
mse          1857.773818
rmse         43.101900
mae          26.614465
r_sqr        0.310309
st_mse       0.001140
st_rmse      0.033761
st_mae       0.019458
st_r_sqr     0.938681
Name: 0, dtype: float64
=====
=====

```

Model Trained with Learning Rate of 0.6000000000000001

```

-----
Final Training results
mae          6.137523
mse         131.414319
r_sqr        0.959792
rmse        11.463608
Name: 999, dtype: float64
-----

```

```

-----
Final Validation results
val_mae      5.821425
val_mse     132.483197
val_r_sqr    0.956994
val_rmse    11.510135
Name: 999, dtype: float64
-----

```

```

-----
Test Set Results
mse          1153.281994
rmse         33.960006
mae          20.625440
r_sqr        0.555230
st_mse       0.000767
st_rmse      0.027698
st_mae       0.015266
st_r_sqr     0.957287
Name: 0, dtype: float64
=====
=====

```

Model Trained with Learning Rate of 0.7000000000000001

Final Training results

mae 5.179923
mse 104.984511
r_sqr 0.965446
rmse 10.246195
Name: 999, dtype: float64

Final Validation results

val_mae 5.386293
val_mse 120.189784
val_r_sqr 0.949562
val_rmse 10.963110
Name: 999, dtype: float64

Test Set Results

mse 777.561693
rmse 27.884793
mae 14.325743
r_sqr 0.804767
st_mse 0.000993
st_rmse 0.031512
st_mae 0.014685
st_r_sqr 0.950056
Name: 0, dtype: float64

Model Trained with Learning Rate of 0.8

Final Training results

mae 6.580505
mse 152.342796
r_sqr 0.947733
rmse 12.342722
Name: 999, dtype: float64

Final Validation results

```
val_mae      7.899153
val_mse      215.759098
val_r_sqr     0.940280
val_rmse     14.688741
Name: 999, dtype: float64
```


Test Set Results

```
mse          99.381783
rmse         9.969041
mae          5.665376
r_sqr        0.967669
st_mse       0.000953
st_rmse      0.030868
st_mae       0.015169
st_r_sqr     0.937715
Name: 0, dtype: float64
```

=====

Model Trained with Learning Rate of 0.9

Final Training results

```
mae          6.368586
mse         154.791285
r_sqr        0.953307
rmse        12.441515
Name: 999, dtype: float64
```


Final Validation results

```
val_mae      6.191001
val_mse     124.999111
val_r_sqr     0.952075
val_rmse     11.180300
Name: 999, dtype: float64
```


Test Set Results

```
mse          353.851397
rmse         18.810938
mae          10.752358
r_sqr        0.879272
st_mse       0.000580
st_rmse      0.024089
```

```
st_mae      0.012475
st_r_sqr     0.960218
Name: 0, dtype: float64
```

```
=====
=====
```

Model Trained with Learning Rate of 1.0

```
-----
-----
```

Final Training results

```
mae      6.446859
mse     141.742626
r_sqr     0.955973
rmse     11.905571
Name: 999, dtype: float64
```

```
-----
-----
```

Final Validation results

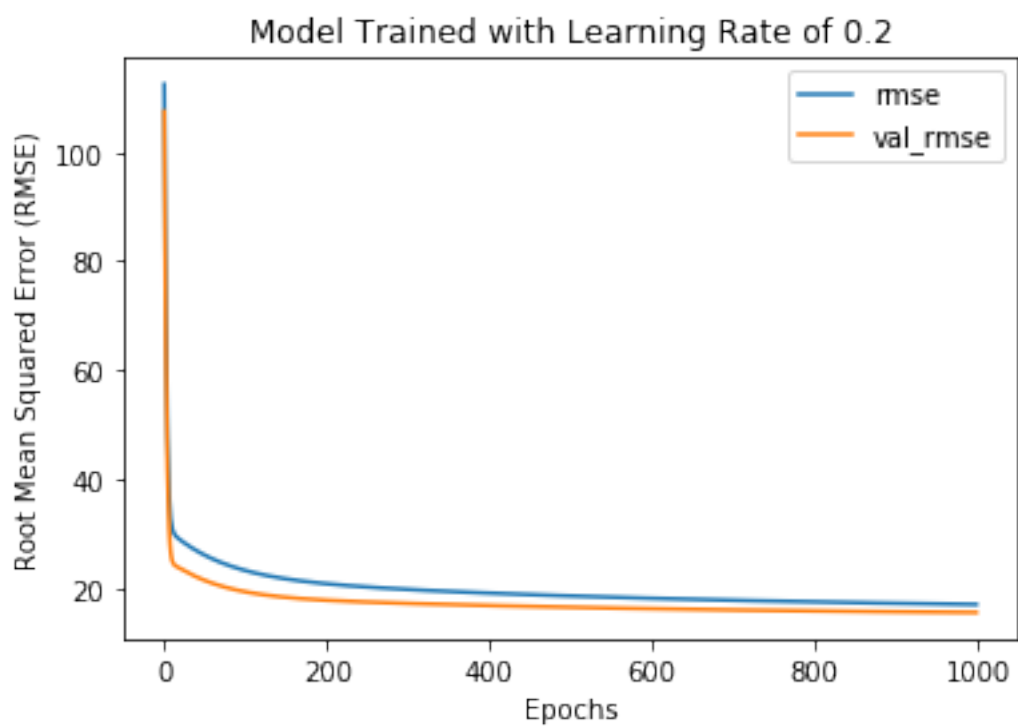
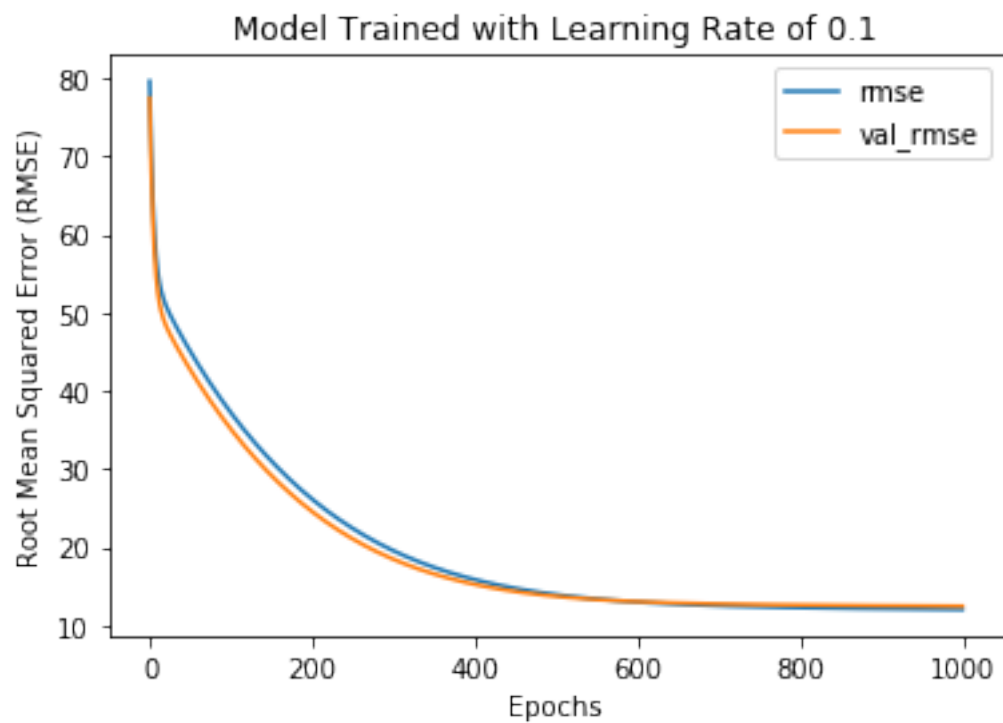
```
val_mae     7.205870
val_mse    183.563748
val_r_sqr    0.938170
val_rmse    13.548570
Name: 999, dtype: float64
```

```
-----
-----
```

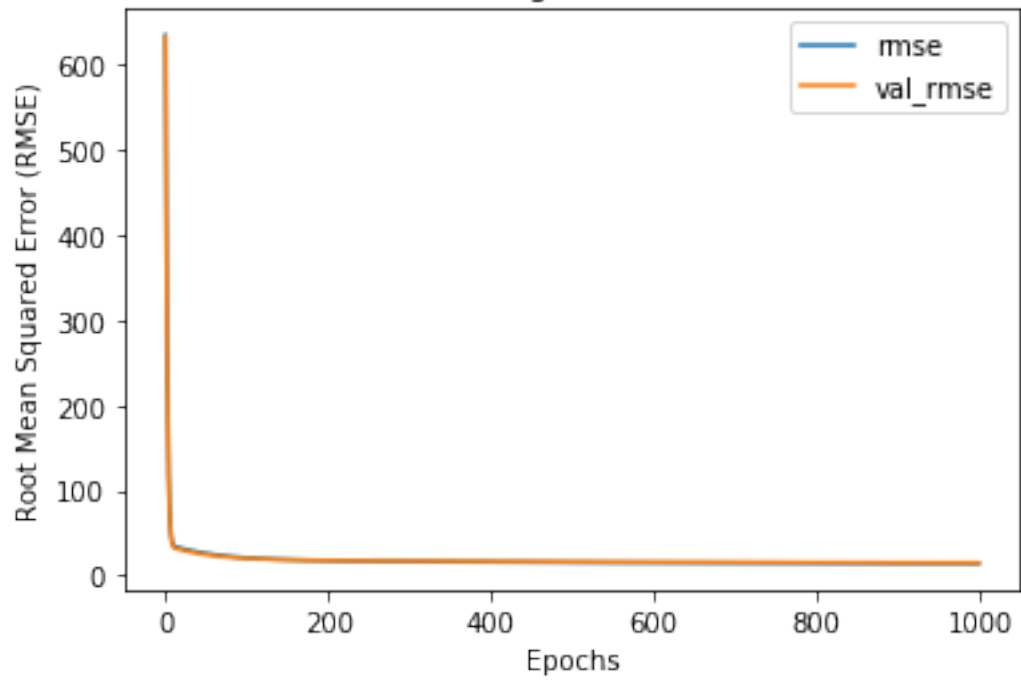
Test Set Results

```
mse      173.410666
rmse     13.168548
mae       7.166303
r_sqr     0.939510
st_mse     0.000876
st_rmse    0.029593
st_mae     0.015108
st_r_sqr    0.938673
Name: 0, dtype: float64
```

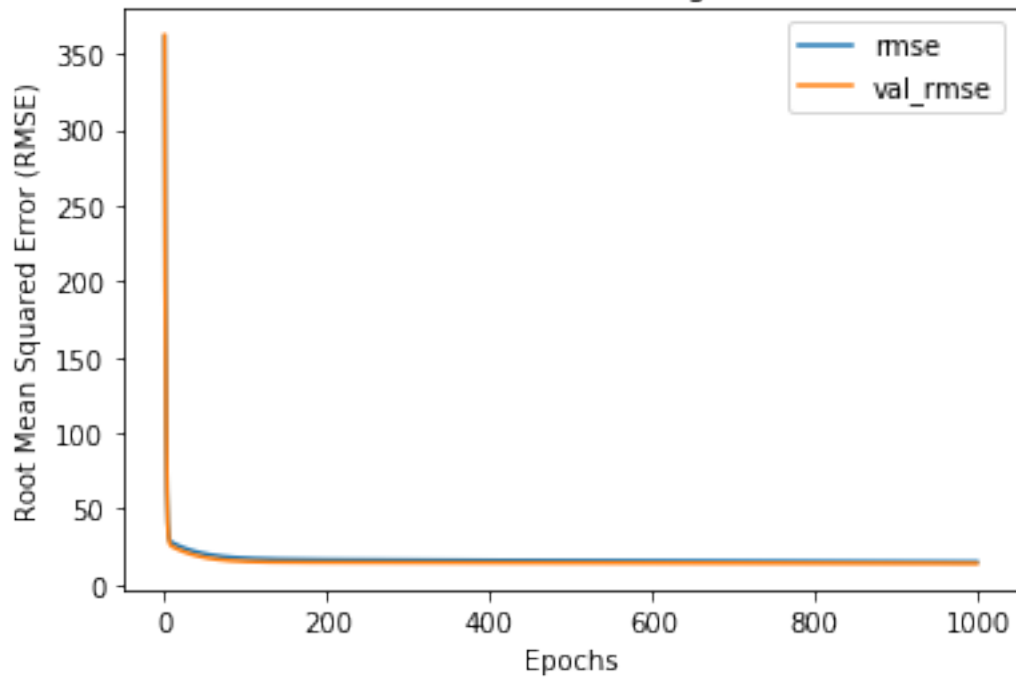
```
=====
=====
```

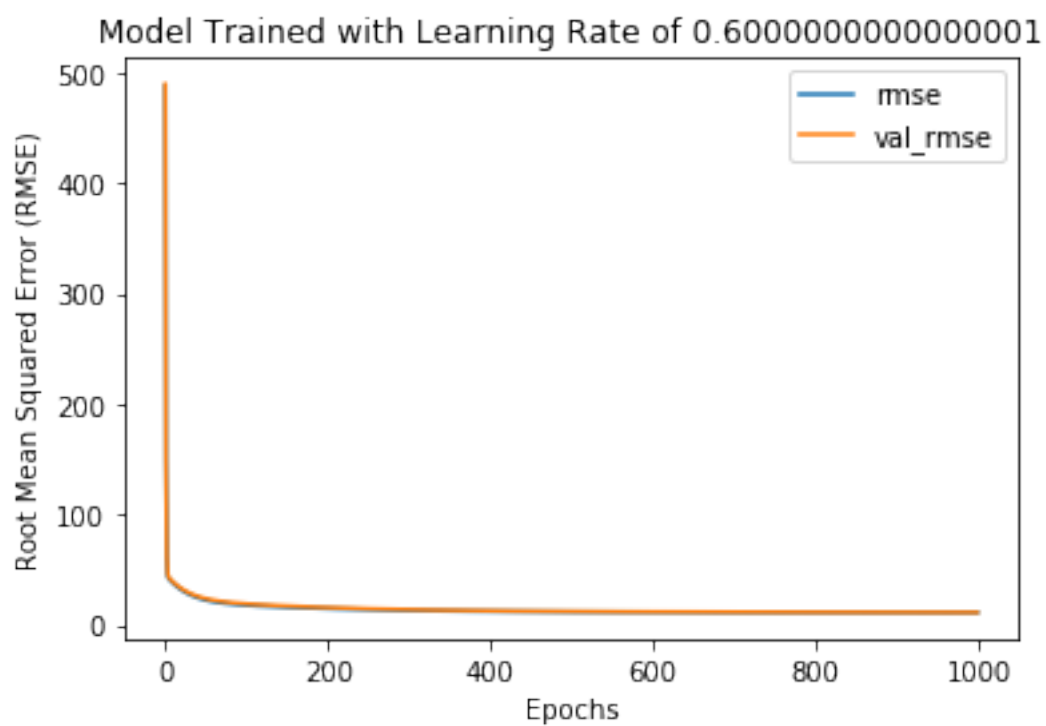
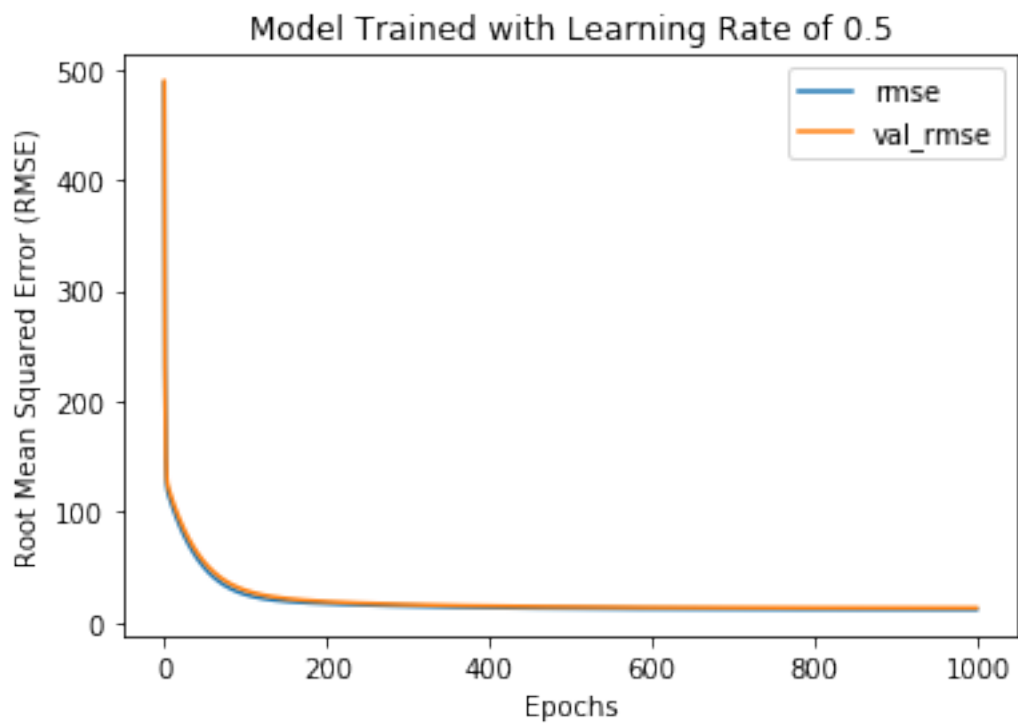


Model Trained with Learning Rate of 0.30000000000000004

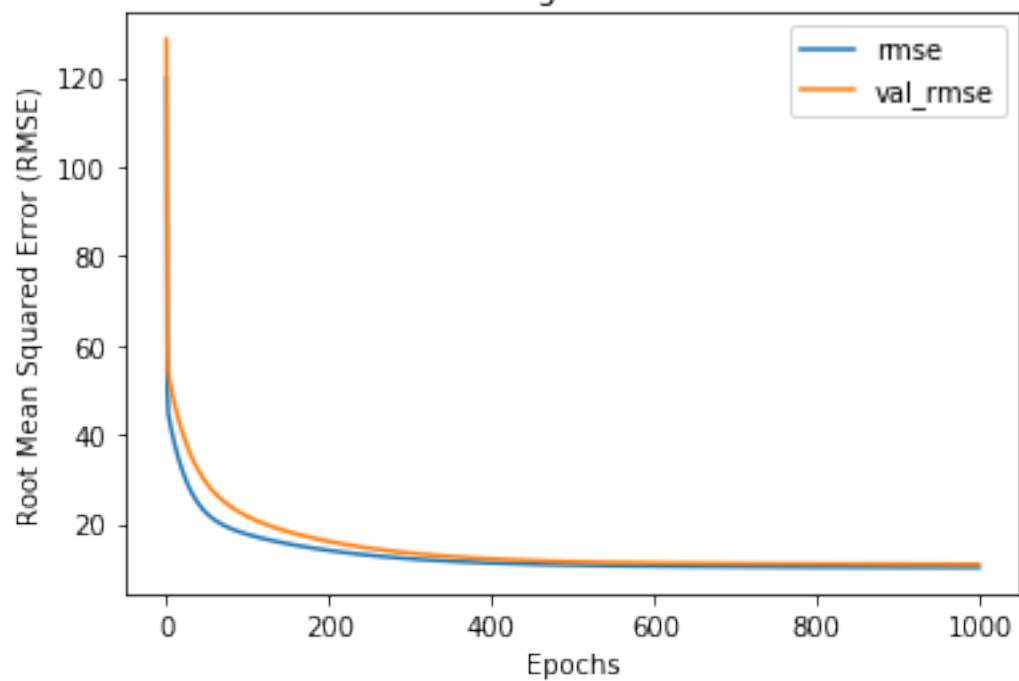


Model Trained with Learning Rate of 0.4

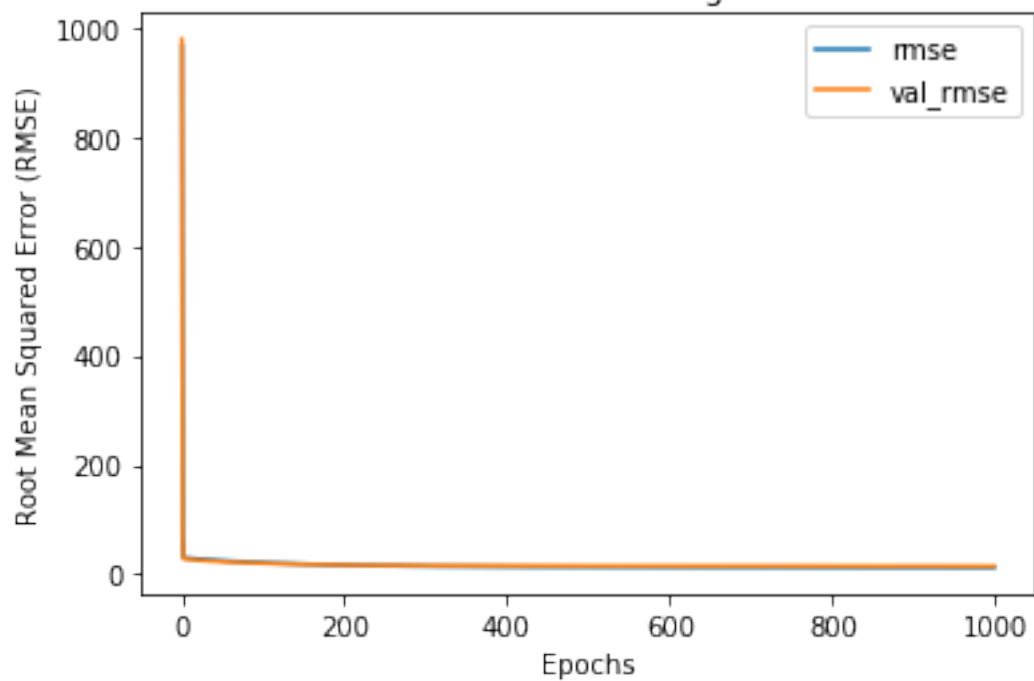


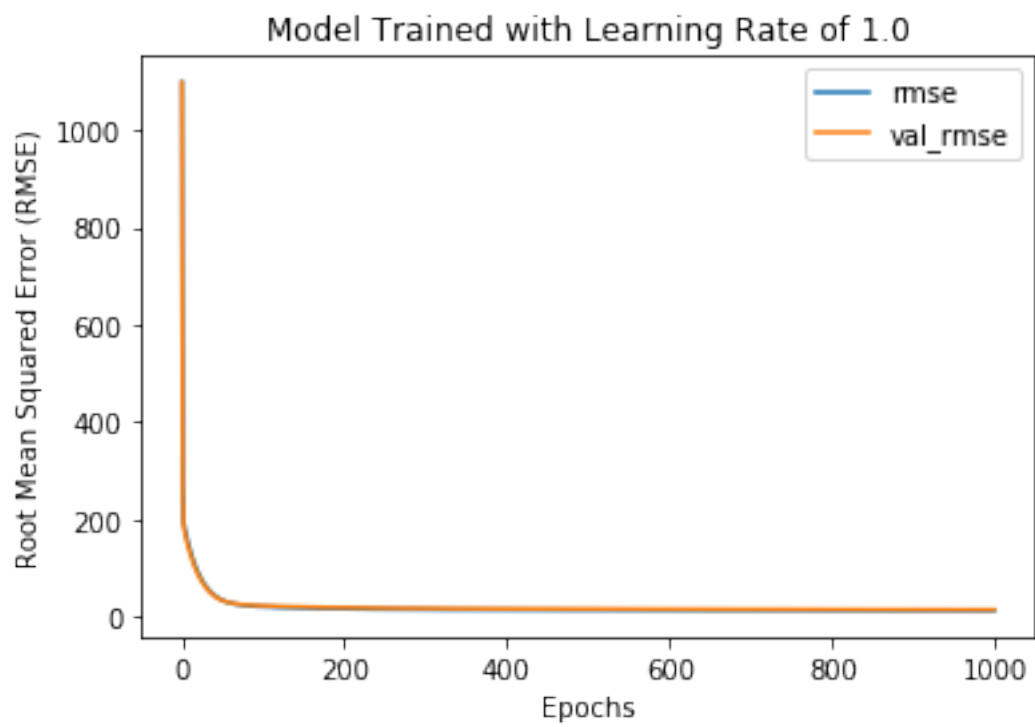
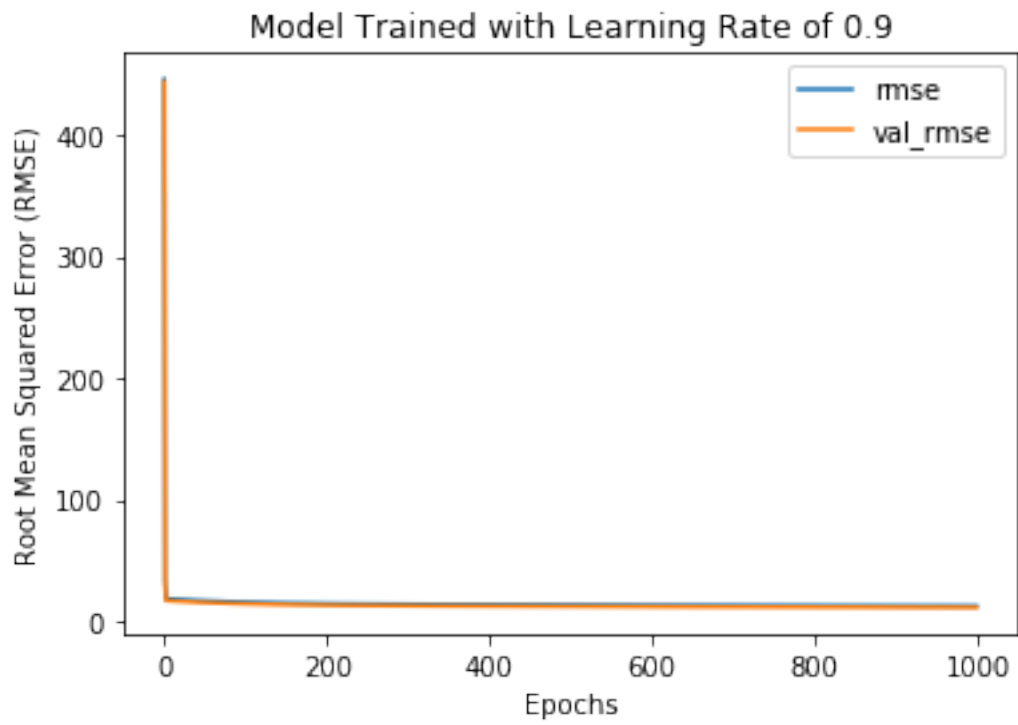


Model Trained with Learning Rate of 0.7000000000000001



Model Trained with Learning Rate of 0.8





- learning rate of 0.3 appears to offer most consistent results between training, validation and testing sets

Activation Functions

```
[15]: activ_tests = dict()
      activations = ("sigmoid", "tanh", "relu", "linear")

      for func in activations:
          activ_tests[func] = build_train_test(
              fs,
              feature_cols,
              target_cols,
              layers=("auto", 1),
              activ_func=func,
              epochs=1000,
              l_rate=0.3
          )

[16]: for a in activations:
      print(f"Model Trained with {a.capitalize()} Activation Function",
        ↪end=f"\n{'-'*100}\n")

      print("Final Training Results")
      print(activ_tests[a]["training_results"].iloc[-1, :4], end=f"\n{'-'*100}\n")

      print("Final Validation Results")
      print(activ_tests[a]["training_results"].iloc[-1, 8:12],
        ↪end=f"\n{'-'*100}\n")

      print("Test Set Results")
      print(activ_tests[a]["error_metrics"].iloc[0], end=f"\n{'-'*100}\n\n\n\n")

      ax = activ_tests[a]["training_results"].plot(
          y=["rmse", "val_rmse"], title=f"Model Trained with {a.capitalize()}
        ↪Activation Function",
      )
      ax.set_xlabel("Epochs")
      ax.set_ylabel("Root Mean Squared Error (RMSE)")
```

Model Trained with Sigmoid Activation Function

```
-----
-----
Final Training Results
mae      38.514325
mse      2535.839810
r_sqr     0.099437
rmse      50.357123
Name: 999, dtype: float64
```

```
-----  
-----  
Final Validation Results  
val_mae      40.178245  
val_mse      3115.092457  
val_r_sqr     0.116448  
val_rmse     55.813013  
Name: 999, dtype: float64  
-----
```

```
-----  
Test Set Results  
mse          3205.065470  
rmse         56.613298  
mae          41.138513  
r_sqr        0.081623  
st_mse       0.010746  
st_rmse      0.103662  
st_mae       0.082285  
st_r_sqr     0.050099  
Name: 0, dtype: float64  
=====
```

```
=====
```

Model Trained with Tanh Activation Function

```
-----
```

```
-----  
Final Training Results  
mae          8.854036  
mse          230.046785  
r_sqr        0.928153  
rmse         15.167293  
Name: 999, dtype: float64  
-----
```

```
-----  
Final Validation Results  
val_mae      7.966150  
val_mse      175.057784  
val_r_sqr     0.927396  
val_rmse     13.230940  
Name: 999, dtype: float64  
-----
```

```
-----  
Test Set Results  
mse          278.640078  
rmse         16.692516  
mae          9.667961
```

```
r_sqr      0.919583
st_mse     0.001965
st_rmse    0.044330
st_mae     0.029884
st_r_sqr   0.825110
Name: 0, dtype: float64
```

```
=====
=====
```

Model Trained with Relu Activation Function

```
-----
-----
```

```
Final Training Results
mae      100.179389
mse     13609.240910
r_sqr    -2.808559
rmse     116.658651
Name: 999, dtype: float64
```

```
-----
-----
```

```
Final Validation Results
val_mae    99.653643
val_mse   12710.176655
val_r_sqr  -3.573111
val_rmse   112.739419
Name: 999, dtype: float64
```

```
-----
-----
```

```
Test Set Results
mse      10807.651568
rmse     103.959856
mae       93.997537
r_sqr    -4.480235
st_mse    0.053553
st_rmse   0.231416
st_mae    0.199983
st_r_sqr  -2.949367
Name: 0, dtype: float64
```

```
=====
=====
```

Model Trained with Linear Activation Function

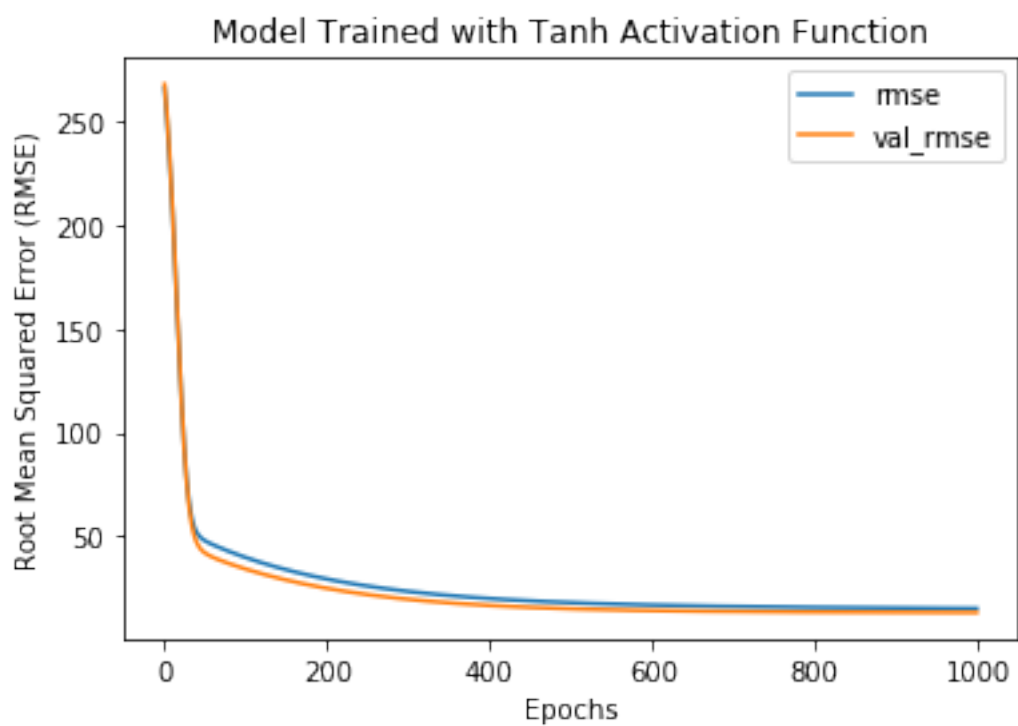
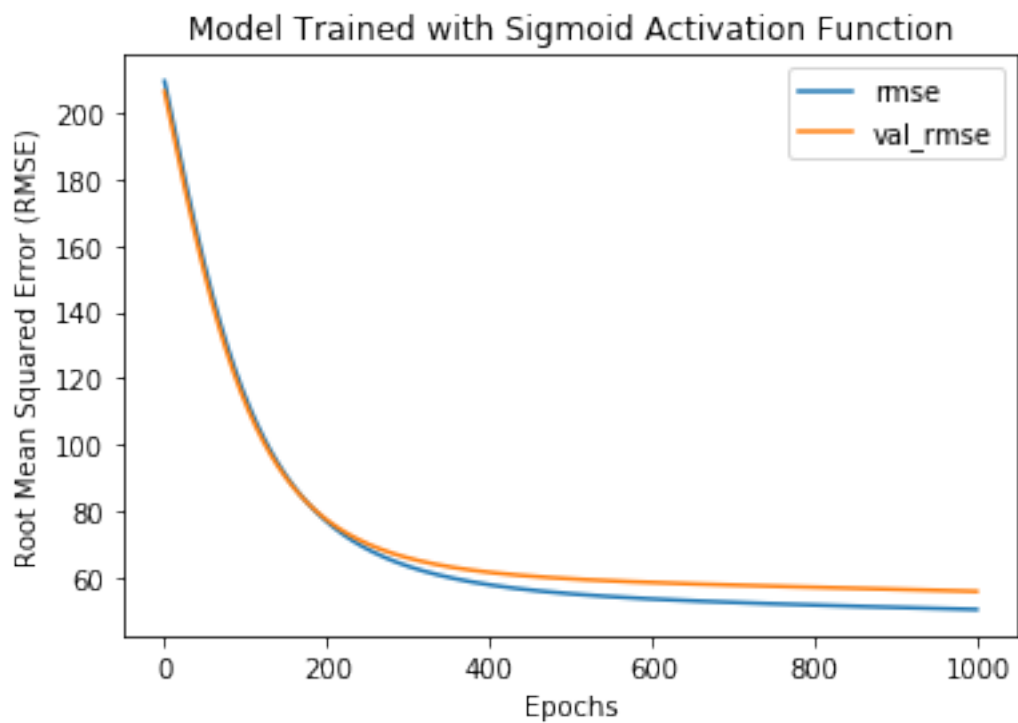
```
-----
-----
```

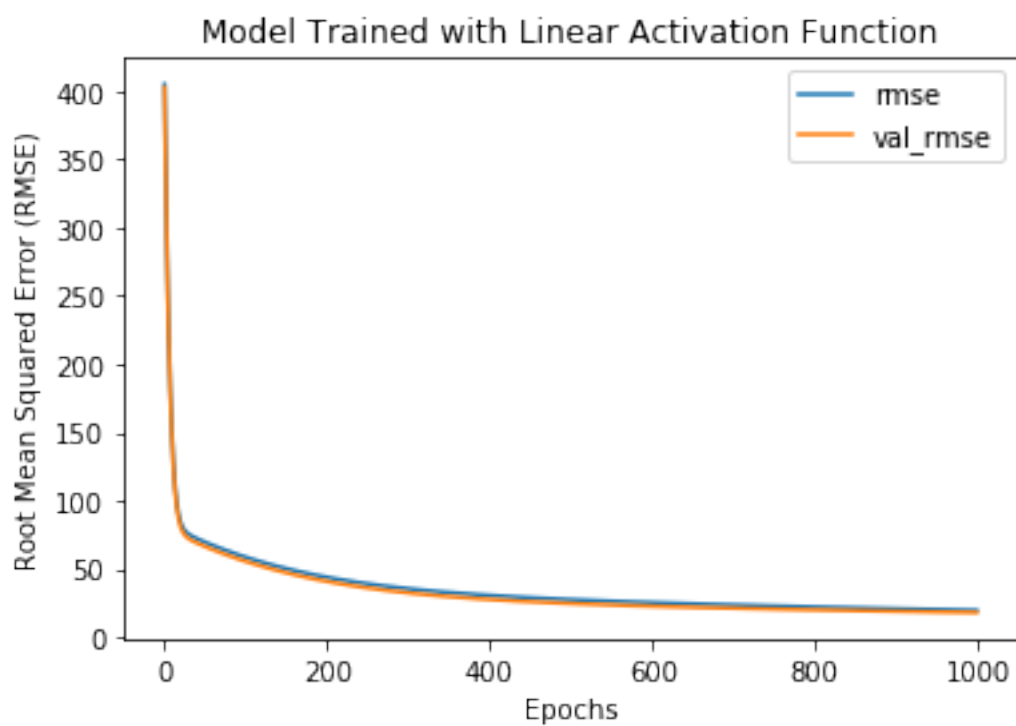
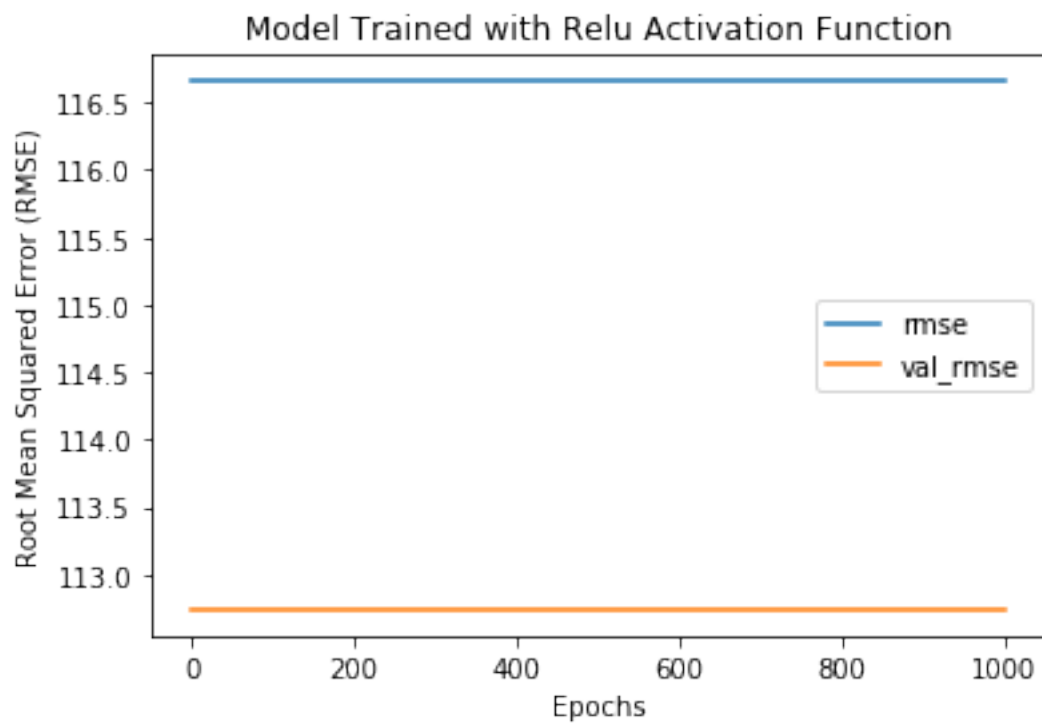
Final Training Results
mae 12.242387
mse 389.890987
r_sqr 0.868005
rmse 19.745657
Name: 999, dtype: float64

Final Validation Results
val_mae 11.112784
val_mse 330.478029
val_r_sqr 0.895496
val_rmse 18.179055
Name: 999, dtype: float64

Test Set Results
mse 496.884198
rmse 22.290899
mae 13.422733
r_sqr 0.856823
st_mse 0.001640
st_rmse 0.040497
st_mae 0.023602
st_r_sqr 0.854173
Name: 0, dtype: float64

=====





- tanh is quite clearly the best performing activation function

Hidden Layers

```
[23]: layer_tests = dict()
for i in range(1, 11):
    layer_tests[f"Test-{i}"] = build_train_test(
        fs,
        feature_cols,
        target_cols,
        layers=("auto", i, 1),
        activ_func="tanh",
        epochs=1000,
        l_rate=0.3
    )
```

```
[24]: for i in range(1,11):
        print(f"Model trained with Single Hidden Layer of Size {i}",
        ↪end=f"\n{'-'*100}\n")

        print("Final Training Results")
        print(layer_tests[f"Test-{i}"]["training_results"].iloc[-1, :4],
        ↪end=f"\n{'-'*100}\n")

        print("Final Validation Results")
        print(layer_tests[f"Test-{i}"]["training_results"].iloc[-1, 8:12],
        ↪end=f"\n{'-'*100}\n")

        print("Test Set Results")
        print(layer_tests[f"Test-{i}"]["error_metrics"].iloc[0],
        ↪end=f"\n{'-'*100}\n\n\n\n")

        ax = layer_tests[f"Test-{i}"]["training_results"].plot(
            y=["rmse", "val_rmse"], title=f"Model trained with Single Hidden Layer
        ↪of Size {i}",
        )
        ax.set_xlabel("Epochs")
        ax.set_ylabel("Root Mean Squared Error (RMSE)")
```

Model trained with Single Hidden Layer of Size 1

Final Training Results

mae	6.537448
mse	141.683314
r_sqr	0.953481
rmse	11.903080

Name: 999, dtype: float64

Final Validation Results

val_mae 6.708576

val_mse 143.885159

val_r_sqr 0.946122

val_rmse 11.995214

Name: 999, dtype: float64

Test Set Results

mse 727.067186

rmse 26.964183

mae 15.432374

r_sqr 0.802375

st_mse 0.001012

st_rmse 0.031812

st_mae 0.017118

st_r_sqr 0.957242

Name: 0, dtype: float64

=====

Model trained with Single Hidden Layer of Size 2

Final Training Results

mae 10.205298

mse 298.325655

r_sqr 0.904494

rmse 17.272106

Name: 999, dtype: float64

Final Validation Results

val_mae 9.514489

val_mse 230.076913

val_r_sqr 0.907843

val_rmse 15.168286

Name: 999, dtype: float64

Test Set Results

mse 510.334071

rmse 22.590575

```
mae          13.883413
r_sqr         0.858922
st_mse        0.003538
st_rmse       0.059479
st_mae        0.027689
st_r_sqr      0.849668
Name: 0, dtype: float64
```

```
=====
=====
```

Model trained with Single Hidden Layer of Size 3

```
-----
-----
```

Final Training Results

```
mae          7.705780
mse         184.419617
r_sqr        0.939006
rmse        13.580118
Name: 999, dtype: float64
```

```
-----
-----
```

Final Validation Results

```
val_mae       8.703655
val_mse       217.514316
val_r_sqr     0.925730
val_rmse      14.748367
Name: 999, dtype: float64
```

```
-----
-----
```

Test Set Results

```
mse          865.644607
rmse         29.421839
mae          20.407450
r_sqr        0.752135
st_mse       0.001487
st_rmse      0.038567
st_mae       0.021021
st_r_sqr     0.933960
Name: 0, dtype: float64
```

```
=====
=====
```

Model trained with Single Hidden Layer of Size 4

```
-----
```

```
-----  
Final Training Results  
mae      11.749247  
mse      349.502596  
r_sqr     0.892044  
rmse     18.694989  
Name: 999, dtype: float64  
-----
```

```
-----  
Final Validation Results  
val_mae    9.944671  
val_mse   190.068800  
val_r_sqr   0.905531  
val_rmse   13.786544  
Name: 999, dtype: float64  
-----
```

```
-----  
Test Set Results  
mse      428.229721  
rmse     20.693712  
mae      14.541870  
r_sqr     0.885195  
st_mse     0.002819  
st_rmse    0.053096  
st_mae     0.028576  
st_r_sqr   0.868651  
Name: 0, dtype: float64  
=====
```

```
=====
```

Model trained with Single Hidden Layer of Size 5

```
-----
```

```
-----  
Final Training Results  
mae      10.956204  
mse      323.253611  
r_sqr     0.891739  
rmse     17.979255  
Name: 999, dtype: float64  
-----
```

```
-----  
Final Validation Results  
val_mae    11.432056  
val_mse   330.366397  
val_r_sqr   0.903258  
val_rmse   18.175984
```

Name: 999, dtype: float64

Test Set Results

mse	583.194248
rmse	24.149415
mae	15.859290
r_sqr	0.812844
st_mse	0.002409
st_rmse	0.049082
st_mae	0.025178
st_r_sqr	0.881684

Name: 0, dtype: float64

=====

Model trained with Single Hidden Layer of Size 6

Final Training Results

mae	8.474634
mse	217.399014
r_sqr	0.917411
rmse	14.744457

Name: 999, dtype: float64

Final Validation Results

val_mae	9.297501
val_mse	301.599923
val_r_sqr	0.904126
val_rmse	17.366632

Name: 999, dtype: float64

Test Set Results

mse	427.127333
rmse	20.667059
mae	10.664607
r_sqr	0.903359
st_mse	0.001479
st_rmse	0.038454
st_mae	0.023261
st_r_sqr	0.896758

Name: 0, dtype: float64

=====

Model trained with Single Hidden Layer of Size 7

Final Training Results

mae 7.878499
mse 185.417026
r_sqr 0.942938
rmse 13.616792
Name: 999, dtype: float64

Final Validation Results

val_mae 8.807907
val_mse 220.582351
val_r_sqr 0.922844
val_rmse 14.852015
Name: 999, dtype: float64

Test Set Results

mse 644.390302
rmse 25.384844
mae 15.620218
r_sqr 0.776858
st_mse 0.002876
st_rmse 0.053625
st_mae 0.027910
st_r_sqr 0.880527
Name: 0, dtype: float64

=====

=====

Model trained with Single Hidden Layer of Size 8

Final Training Results

mae 9.326841
mse 261.368740
r_sqr 0.922007
rmse 16.166903
Name: 999, dtype: float64

```
-----  
Final Validation Results  
val_mae      9.086521  
val_mse     220.420668  
val_r_sqr    0.929665  
val_rmse     14.846571  
Name: 999, dtype: float64  
-----
```

```
-----  
Test Set Results  
mse          459.131447  
rmse         21.427353  
mae          12.533972  
r_sqr        0.799316  
st_mse       0.003604  
st_rmse      0.060036  
st_mae       0.032903  
st_r_sqr     0.816063  
Name: 0, dtype: float64  
=====
```

Model trained with Single Hidden Layer of Size 9

```
-----  
Final Training Results  
mae          8.846222  
mse         208.756315  
r_sqr        0.933647  
rmse        14.448402  
Name: 999, dtype: float64  
-----
```

```
-----  
Final Validation Results  
val_mae      8.027564  
val_mse     162.109535  
val_r_sqr    0.940558  
val_rmse     12.732224  
Name: 999, dtype: float64  
-----
```

```
-----  
Test Set Results  
mse          329.562140  
rmse         18.153846  
mae          12.064198  
r_sqr        0.900492
```

```
st_mse      0.001648
st_rmse     0.040592
st_mae      0.021139
st_r_sqr    0.900036
Name: 0, dtype: float64
```

```
=====
=====
```

Model trained with Single Hidden Layer of Size 10

Final Training Results

```
mae      8.088137
mse     202.259429
r_sqr    0.930396
rmse    14.221794
Name: 999, dtype: float64
```

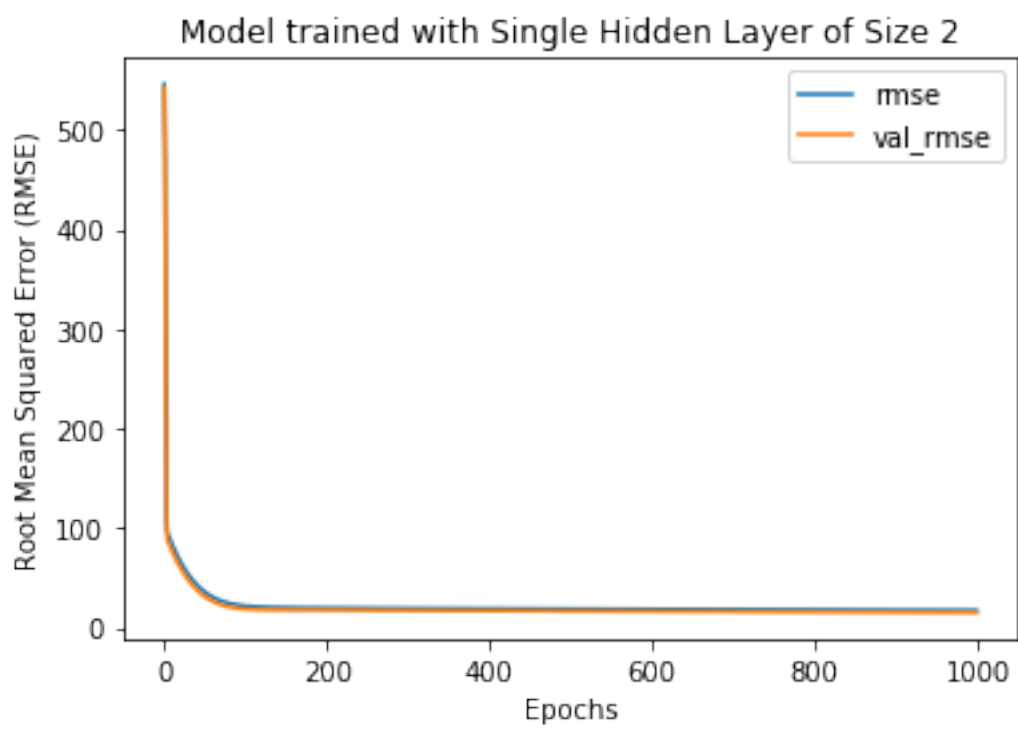
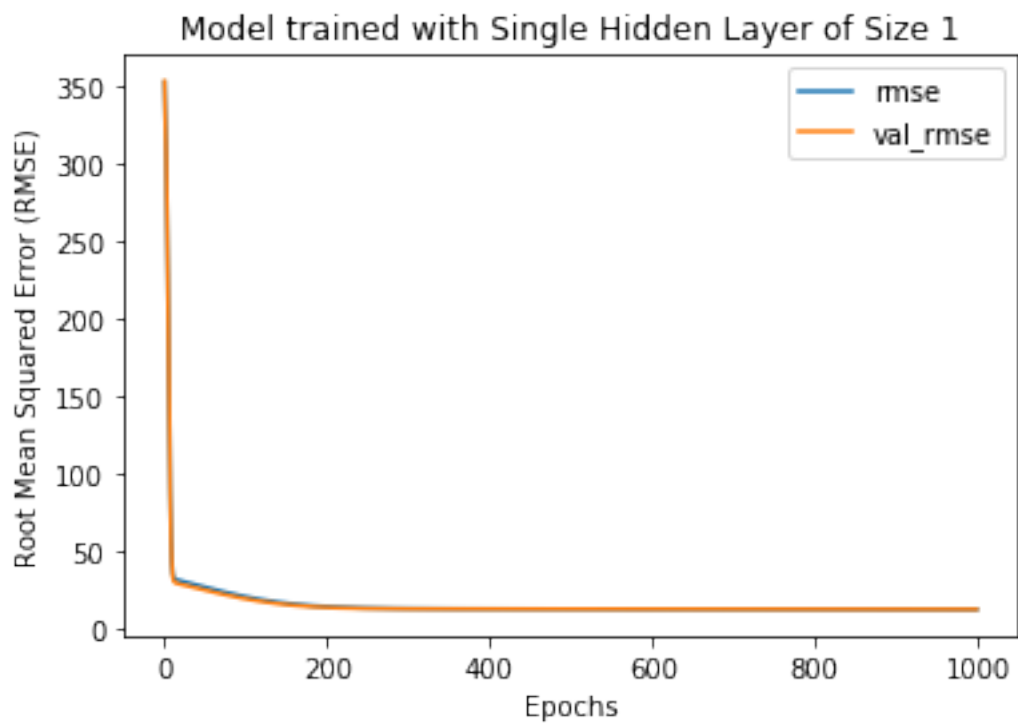
Final Validation Results

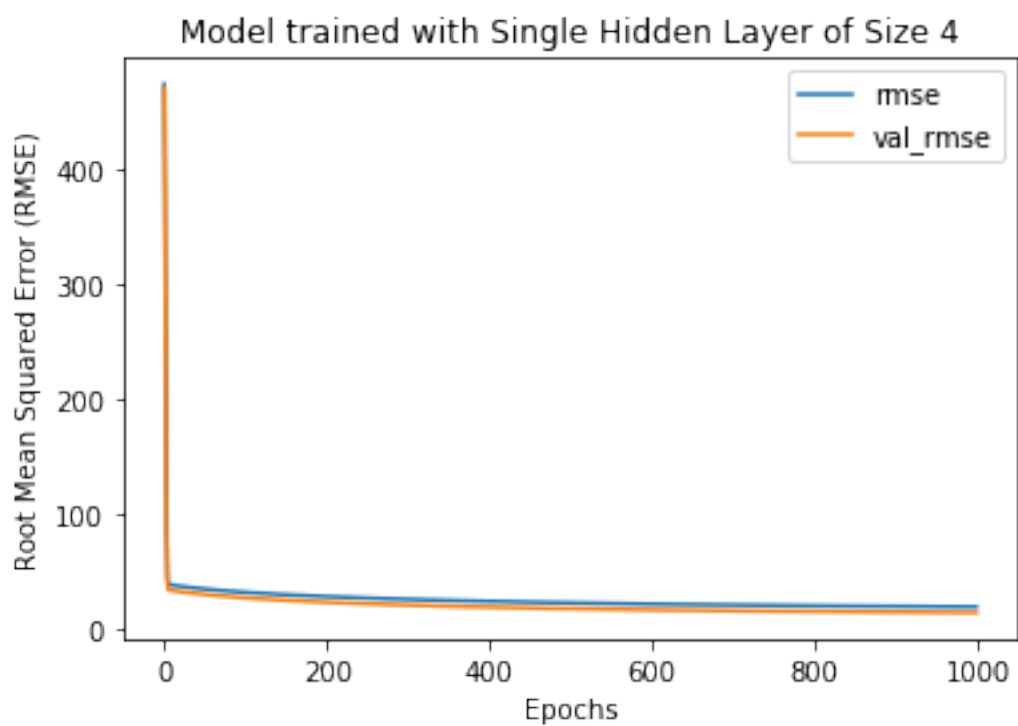
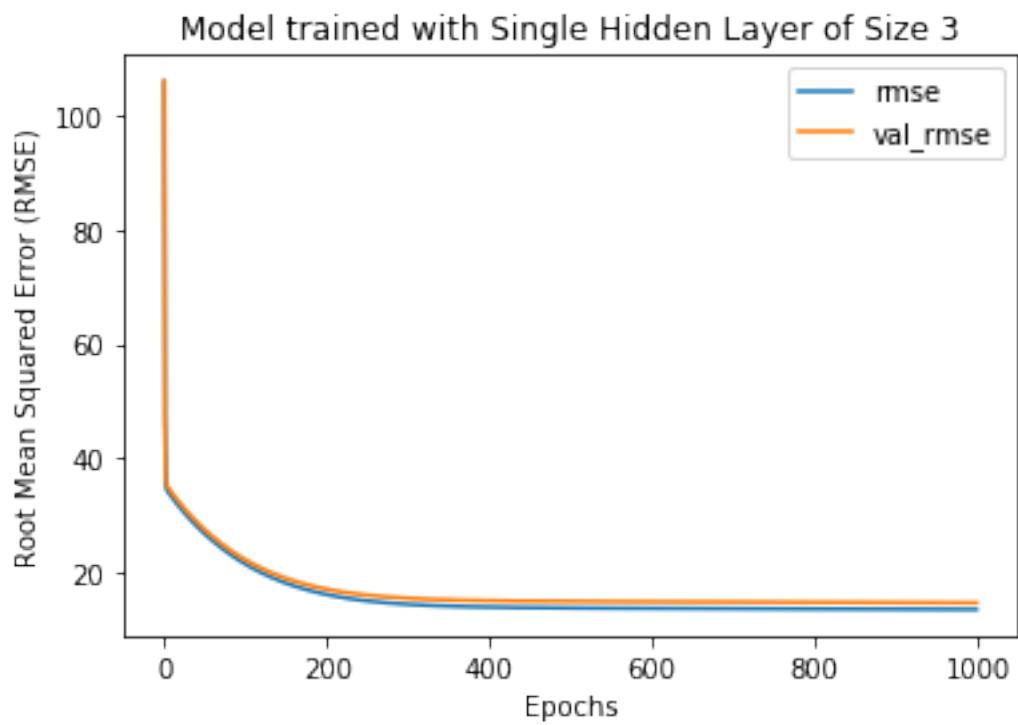
```
val_mae    8.665697
val_mse   191.862063
val_r_sqr   0.951003
val_rmse   13.851428
Name: 999, dtype: float64
```

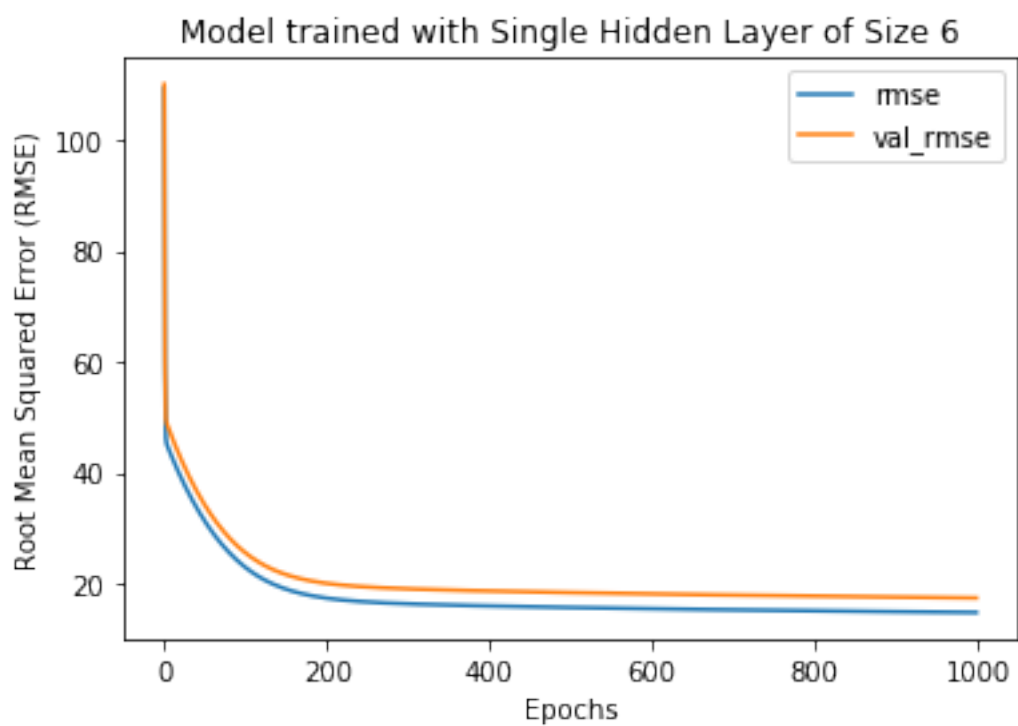
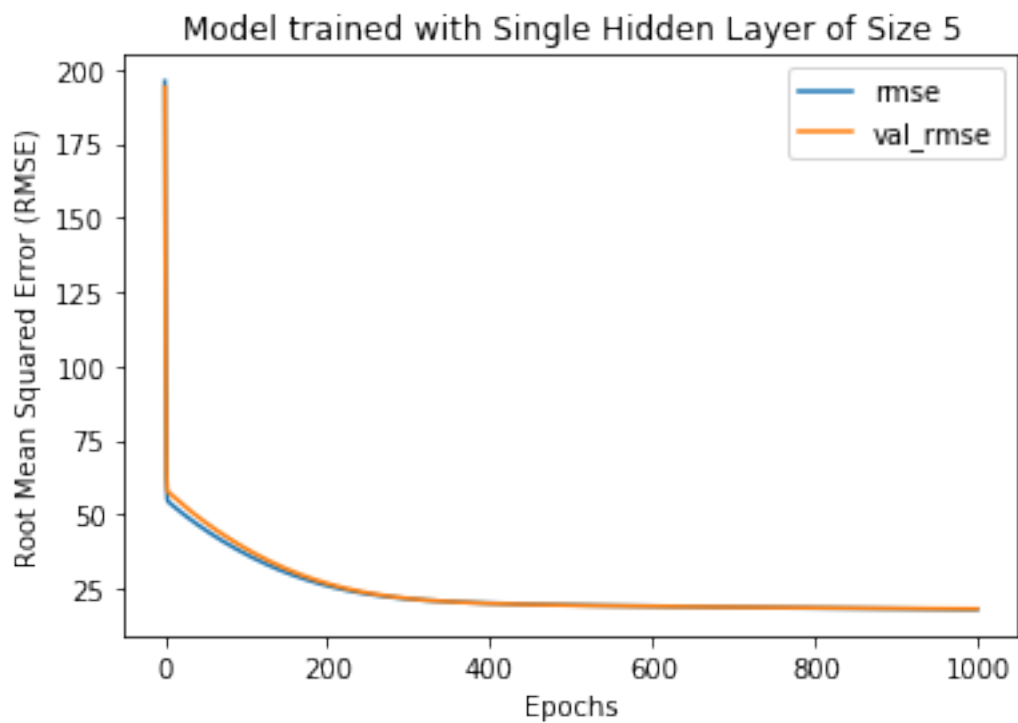
Test Set Results

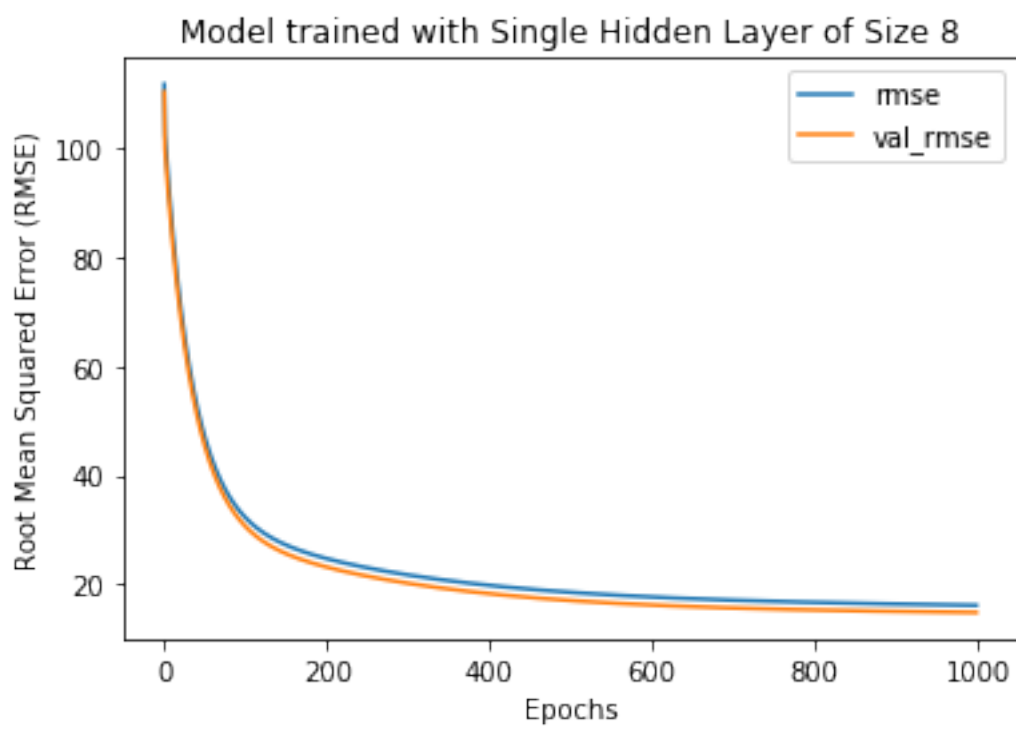
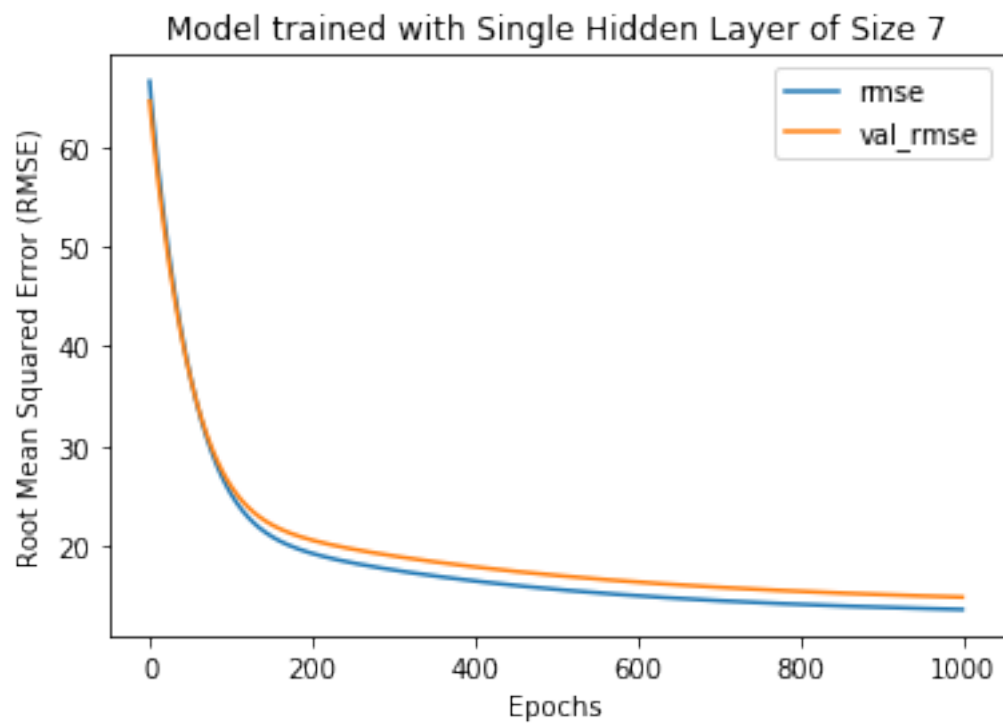
```
mse      178.686358
rmse     13.367362
mae       8.343952
r_sqr     0.936293
st_mse    0.001677
st_rmse   0.040955
st_mae    0.022716
st_r_sqr   0.880053
Name: 0, dtype: float64
```

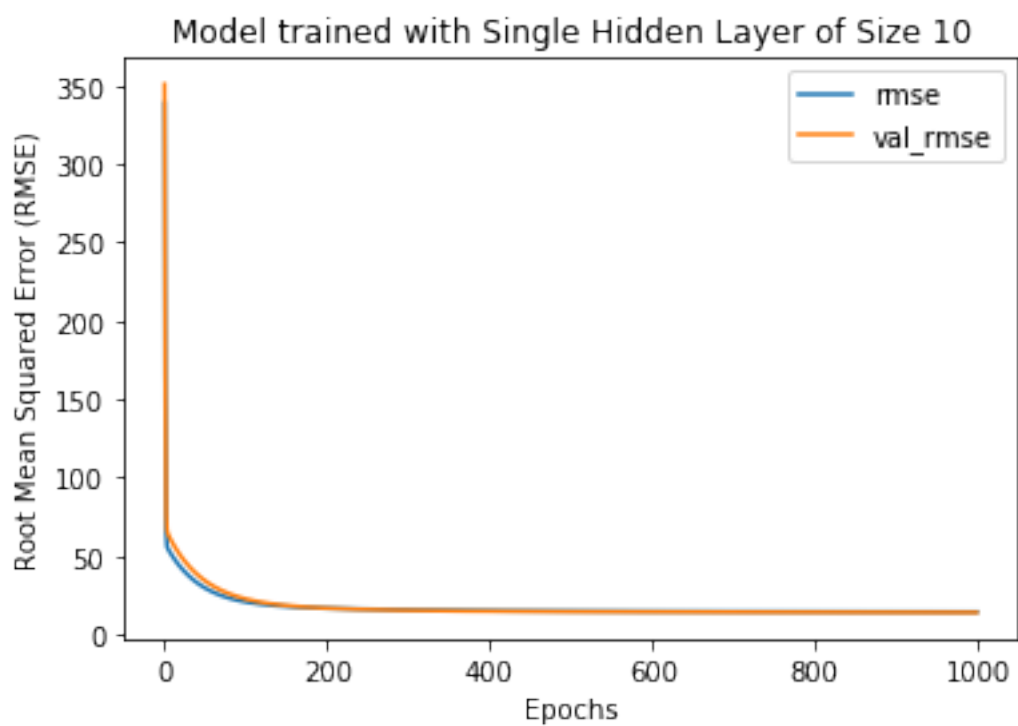
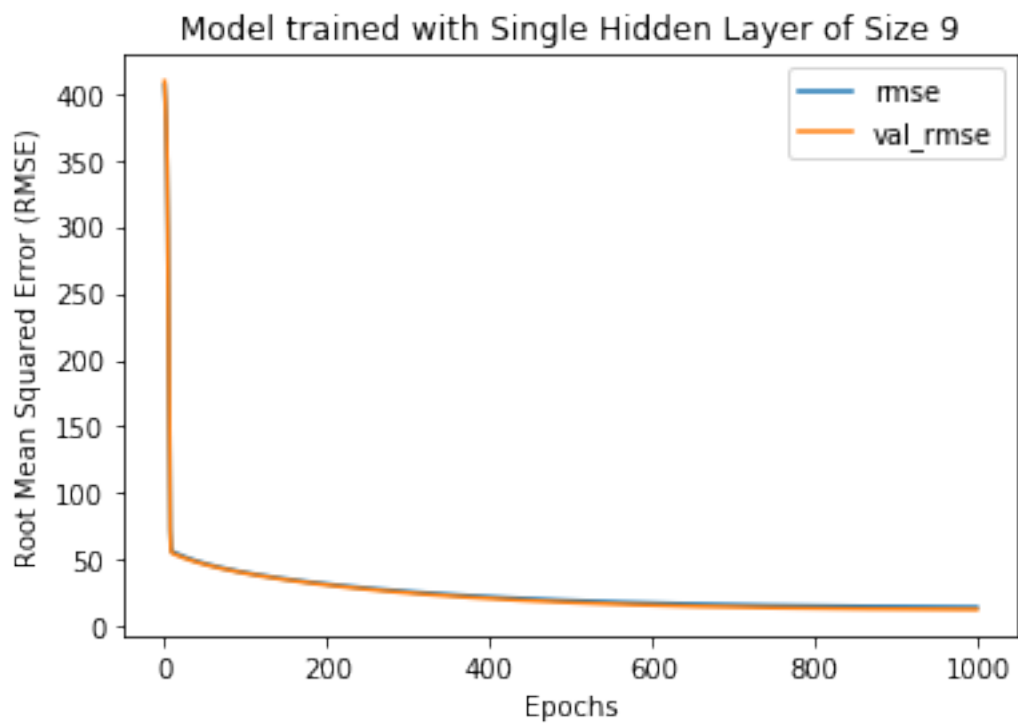
```
=====
=====
```











```
[28]: layer_tests_2 = dict()
for i in range(1, 11):
    layer_tests_2[f"Test-{i}"] = build_train_test(
        fs,
        feature_cols,
        target_cols,
        layers=("auto", 10, i, 1),
        activ_func="tanh",
        epochs=1000,
        l_rate=0.3
    )
```

```
[29]: for i in range(1,11):
    print(f"Model Trained with Hidden Layer Configuration (10, {i})",
    ↪end=f"\n{'-'*100}\n")

    print("Final Training Results")
    print(layer_tests_2[f"Test-{i}"]["training_results"].iloc[-1, :4],
    ↪end=f"\n{'-'*100}\n")

    print("Final Validation Results")
    print(layer_tests_2[f"Test-{i}"]["training_results"].iloc[-1, 8:12],
    ↪end=f"\n{'-'*100}\n")

    print("Test Set Results")
    print(layer_tests_2[f"Test-{i}"]["error_metrics"].iloc[0],
    ↪end=f"\n{'-'*100}\n\n\n")

    ax = layer_tests_2[f"Test-{i}"]["training_results"].plot(
        y=["rmse", "val_rmse"], title=f"Model Trained with Hidden Layer
    ↪Configuration (10, {i})",
    )
    ax.set_xlabel("Epochs")
    ax.set_ylabel("Root Mean Squared Error (RMSE)")
```

Model Trained with Hidden Layer Configuration (10, 1)

Final Training Results

```
mae      11.039224
mse      372.824255
r_sqr    0.892309
rmse     19.308658
Name: 999, dtype: float64
```

Final Validation Results


```
val_mae      8.767553
val_mse     197.875779
val_r_sqr    0.918034
val_rmse     14.066833
Name: 999, dtype: float64
```


Test Set Results

```
mse          327.410711
rmse         18.094494
mae          11.472224
r_sqr        0.877149
st_mse       0.006786
st_rmse      0.082379
st_mae       0.046483
st_r_sqr     0.703177
Name: 0, dtype: float64
```

=====

Model Trained with Hidden Layer Configuration (10, 2)

Final Training Results

```
mae          14.960459
mse          528.579588
r_sqr        0.827110
rmse         22.990859
Name: 999, dtype: float64
```


Final Validation Results

```
val_mae      13.888633
val_mse     505.515482
val_r_sqr    0.845042
val_rmse     22.483671
Name: 999, dtype: float64
```


Test Set Results

```
mse          930.783480
rmse         30.508744
mae          22.436335
r_sqr        0.695578
st_mse       0.003913
st_rmse      0.062555
```

```
st_mae      0.035549
st_r_sqr     0.801065
Name: 0, dtype: float64
```

```
=====
=====
```

Model Trained with Hidden Layer Configuration (10, 3)

```
-----
-----
```

Final Training Results

```
mae      8.733926
mse     228.135558
r_sqr    0.933564
rmse    15.104157
Name: 999, dtype: float64
```

```
-----
-----
```

Final Validation Results

```
val_mae    7.949196
val_mse   168.438061
val_r_sqr   0.932911
val_rmse   12.978369
Name: 999, dtype: float64
```

```
-----
-----
```

Test Set Results

```
mse      1792.981943
rmse      42.343617
mae       29.433090
r_sqr     0.327214
st_mse     0.002311
st_rmse    0.048068
st_mae     0.030563
st_r_sqr   0.865453
Name: 0, dtype: float64
```

```
=====
=====
```

Model Trained with Hidden Layer Configuration (10, 4)

```
-----
-----
```

Final Training Results

```
mae      10.329634
mse     297.055861
```

```
r_sqr      0.904130
rmse       17.235309
Name: 999, dtype: float64
```

Final Validation Results

```
val_mae      10.720939
val_mse      273.030036
val_r_sqr     0.916427
val_rmse     16.523621
Name: 999, dtype: float64
```

Test Set Results

```
mse          297.127479
rmse         17.237386
mae          11.082555
r_sqr        0.898321
st_mse       0.001642
st_rmse      0.040518
st_mae       0.022925
st_r_sqr     0.887155
Name: 0, dtype: float64
```

Model Trained with Hidden Layer Configuration (10, 5)

Final Training Results

```
mae          6.650090
mse         154.343884
r_sqr        0.948926
rmse        12.423521
Name: 999, dtype: float64
```

Final Validation Results

```
val_mae      7.309605
val_mse      178.053204
val_r_sqr     0.945661
val_rmse     13.343658
Name: 999, dtype: float64
```

Test Set Results

```

mse          301.371642
rmse         17.360059
mae          11.063077
r_sqr        0.904170
st_mse       0.001250
st_rmse      0.035358
st_mae       0.016522
st_r_sqr     0.931113
Name: 0, dtype: float64
=====
=====

```

Model Trained with Hidden Layer Configuration (10, 6)

```

-----
Final Training Results
mae          9.374567
mse         238.607777
r_sqr        0.927337
rmse        15.446934
Name: 999, dtype: float64
-----

```

```

-----
Final Validation Results
val_mae       9.233305
val_mse      277.627647
val_r_sqr     0.925936
val_rmse     16.662162
Name: 999, dtype: float64
-----

```

```

-----
Test Set Results
mse          2038.522945
rmse         45.150005
mae          32.770927
r_sqr       -0.095314
st_mse       0.002843
st_rmse      0.053322
st_mae       0.028133
st_r_sqr     0.887762
Name: 0, dtype: float64
=====
=====

```

Model Trained with Hidden Layer Configuration (10, 7)

Final Training Results

```
mae      6.788405
mse     155.399767
r_sqr    0.943998
rmse     12.465944
Name: 999, dtype: float64
```

Final Validation Results

```
val_mae    7.665292
val_mse   240.926042
val_r_sqr   0.936648
val_rmse   15.521792
Name: 999, dtype: float64
```

Test Set Results

```
mse      505.007571
rmse     22.472373
mae      14.885314
r_sqr     0.849342
st_mse     0.001550
st_rmse    0.039370
st_mae     0.019088
st_r_sqr    0.919637
Name: 0, dtype: float64
```

Model Trained with Hidden Layer Configuration (10, 8)

Final Training Results

```
mae      9.593583
mse     264.748185
r_sqr    0.911685
rmse     16.271084
Name: 999, dtype: float64
```

Final Validation Results

```
val_mae    9.723533
val_mse   261.125326
```

```
val_r_sqr      0.918077
val_rmse       16.159373
Name: 999, dtype: float64
```

Test Set Results

```
mse           257.575178
rmse          16.049149
mae           9.836957
r_sqr         0.922273
st_mse        0.002280
st_rmse       0.047754
st_mae        0.024605
st_r_sqr      0.861724
Name: 0, dtype: float64
```

Model Trained with Hidden Layer Configuration (10, 9)

Final Training Results

```
mae           7.164161
mse          176.956075
r_sqr         0.938995
rmse          13.302484
Name: 999, dtype: float64
```

Final Validation Results

```
val_mae       9.147647
val_mse       305.478118
val_r_sqr     0.910433
val_rmse      17.477932
Name: 999, dtype: float64
```

Test Set Results

```
mse           470.744577
rmse          21.696649
mae           12.374243
r_sqr         0.858922
st_mse        0.005060
st_rmse       0.071134
st_mae        0.038185
st_r_sqr      0.818139
```

Name: 0, dtype: float64

=====

Model Trained with Hidden Layer Configuration (10, 10)

Final Training Results

mae 8.182189

mse 216.348942

r_sqr 0.923964

rmse 14.708805

Name: 999, dtype: float64

Final Validation Results

val_mae 9.491528

val_mse 295.626709

val_r_sqr 0.916262

val_rmse 17.193799

Name: 999, dtype: float64

Test Set Results

mse 418.041829

rmse 20.446071

mae 10.285933

r_sqr 0.877130

st_mse 0.000812

st_rmse 0.028500

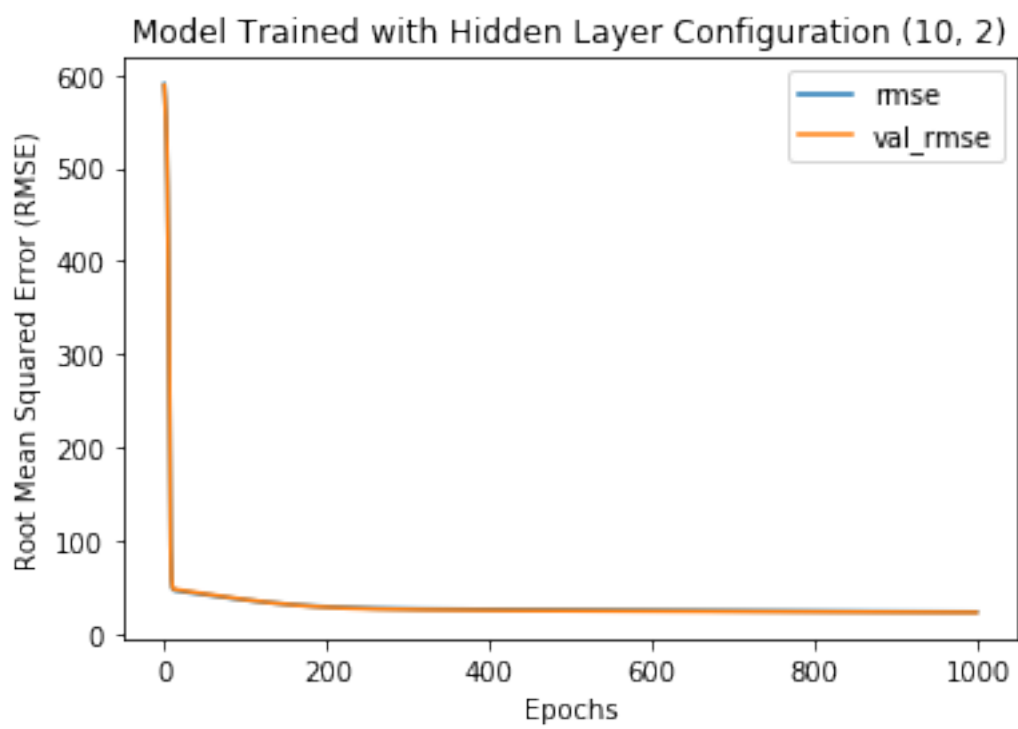
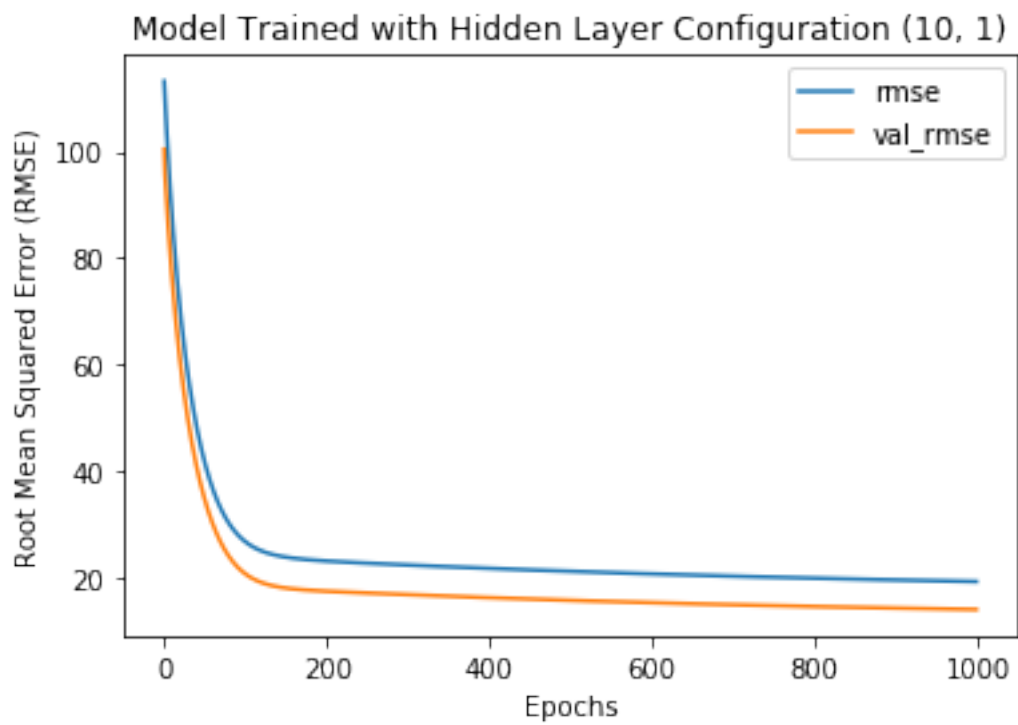
st_mae 0.015443

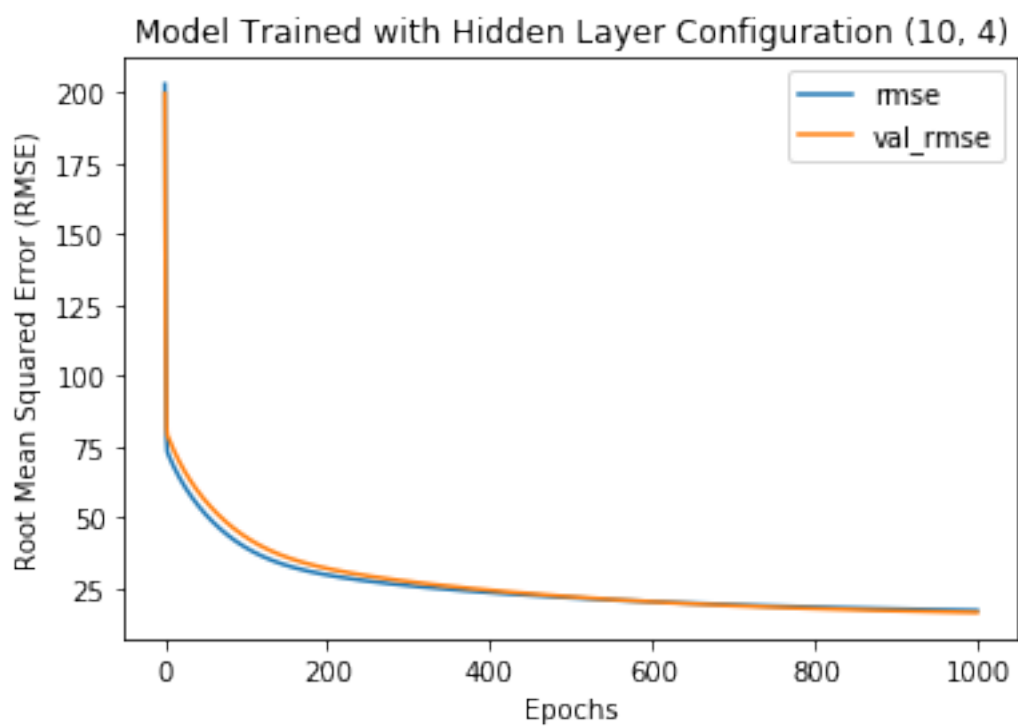
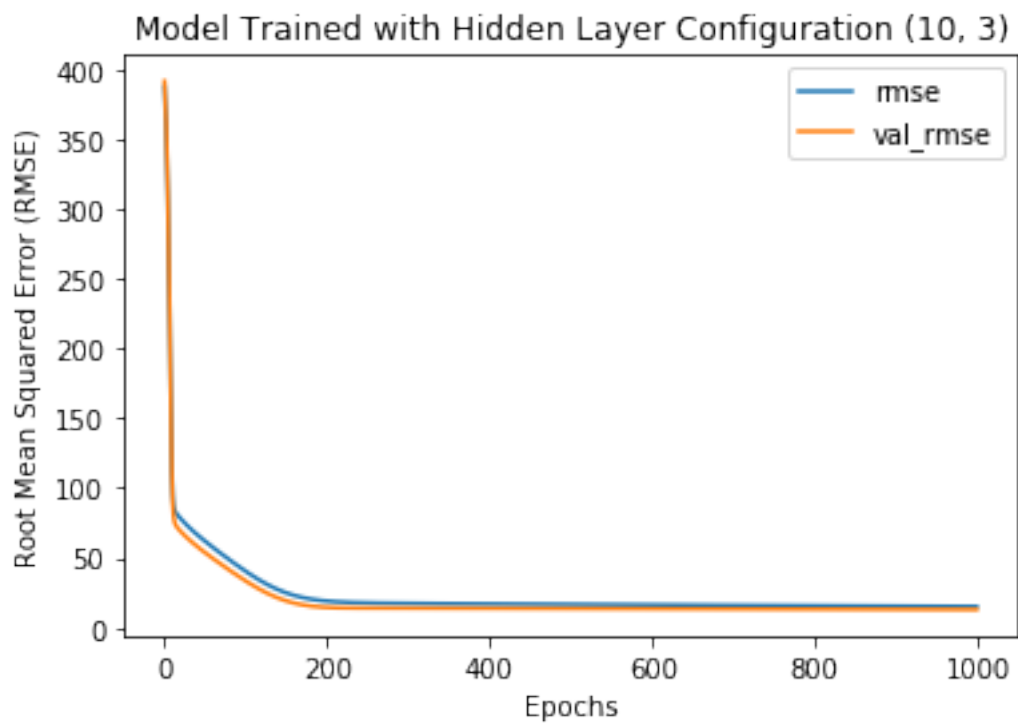
st_r_sqr 0.926349

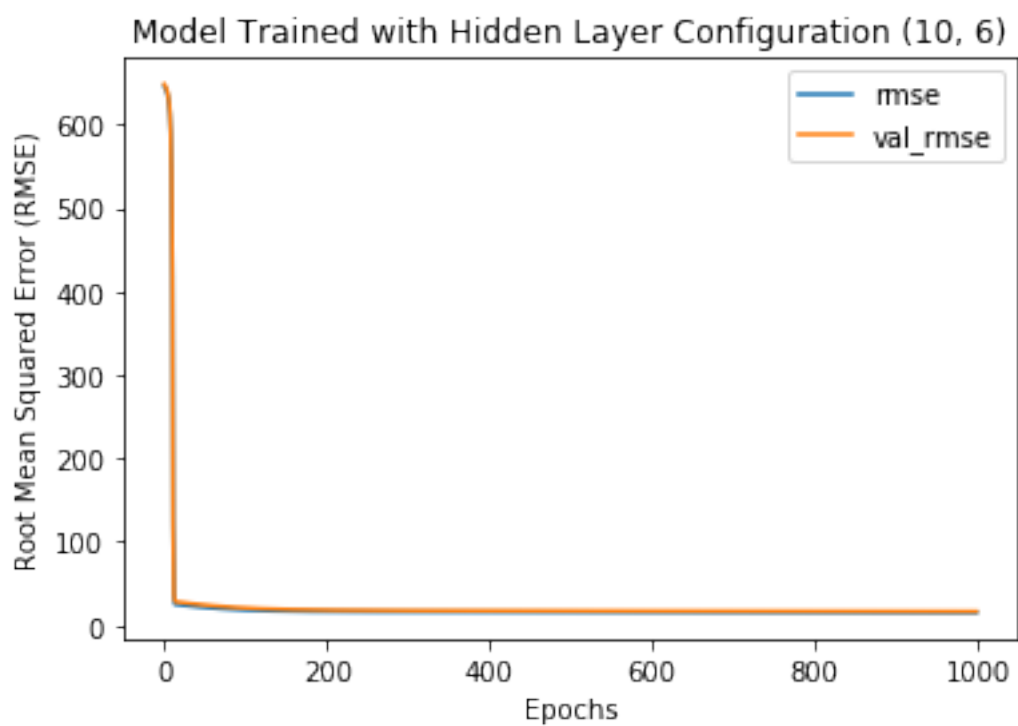
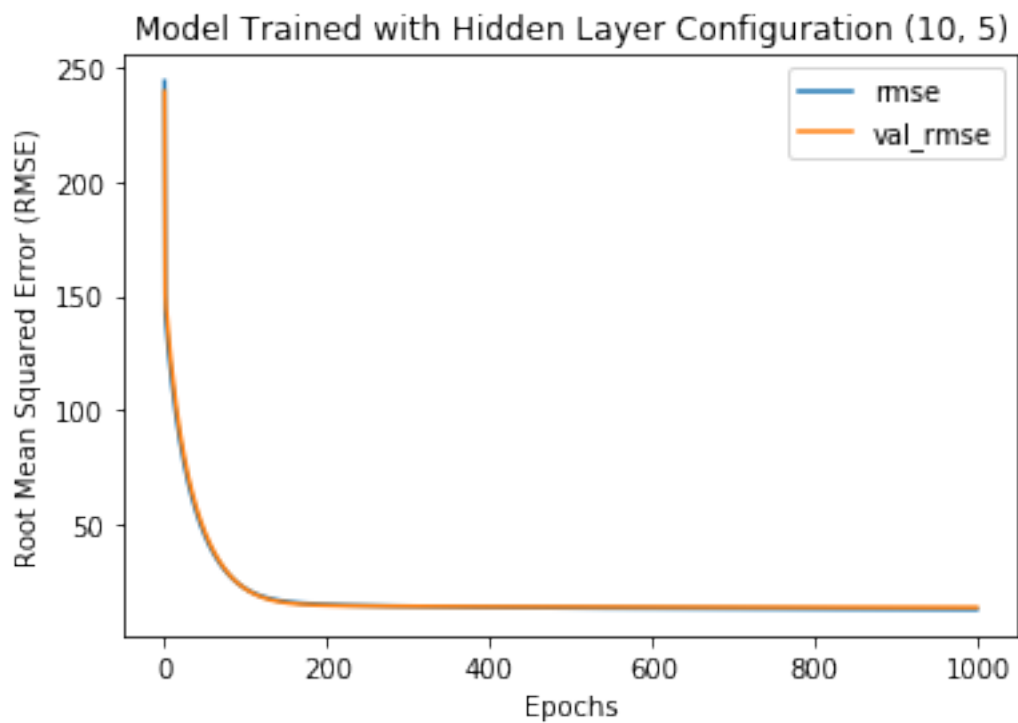
Name: 0, dtype: float64

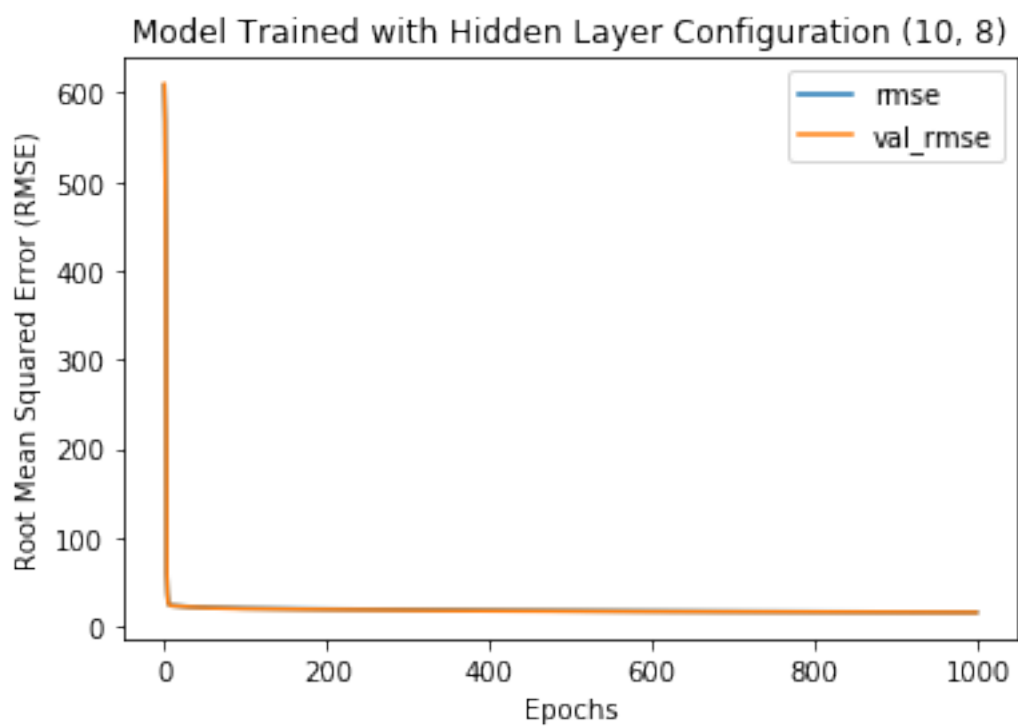
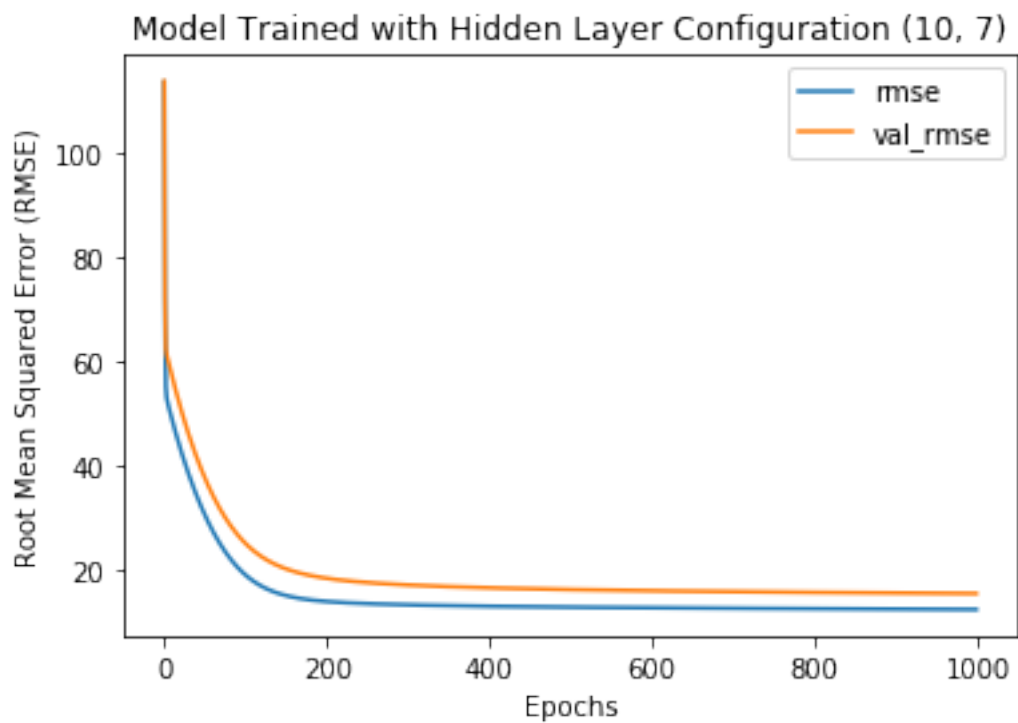
=====

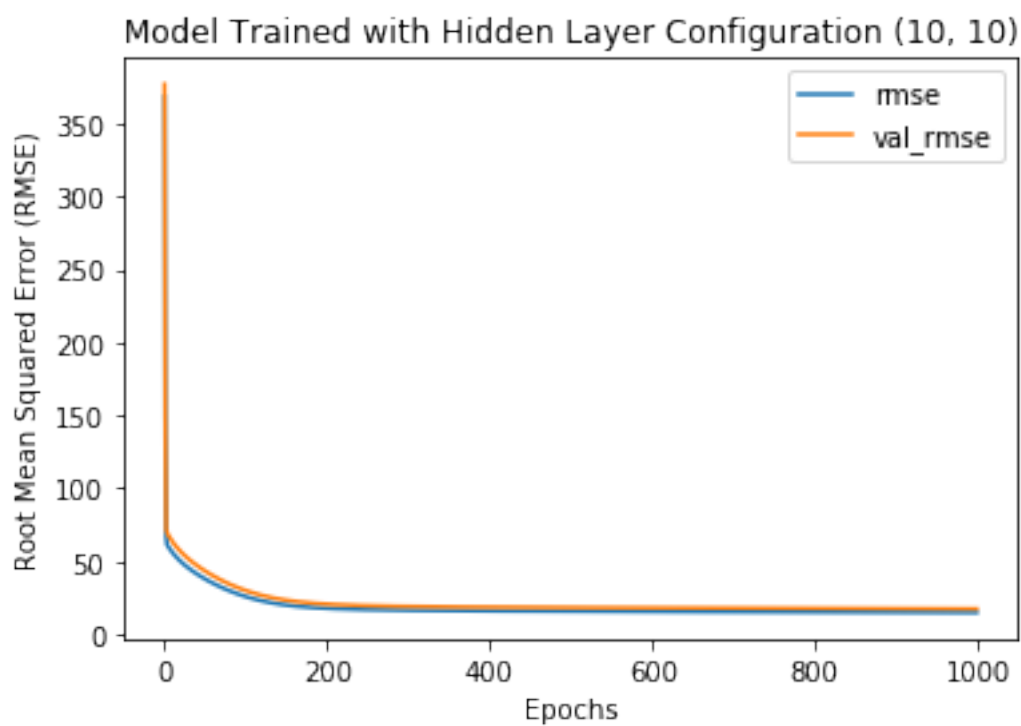
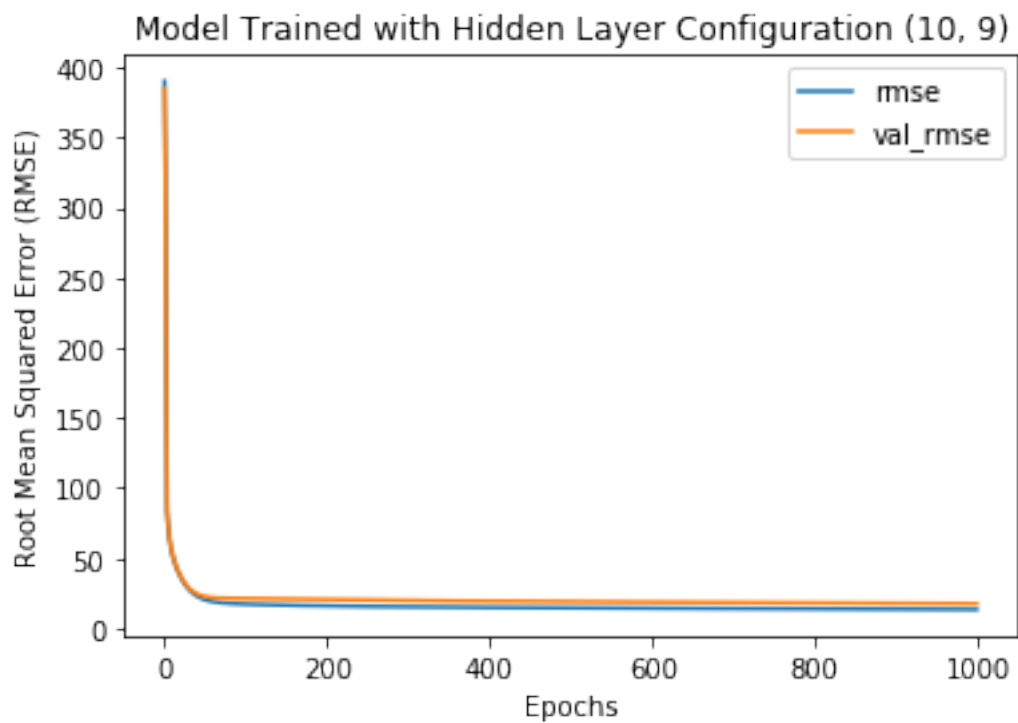
=====











```
[32]: layer_tests_3 = dict()
      for i in range(1, 11):
          layer_tests_3[f"Test-{i}"] = build_train_test(
              fs,
              feature_cols,
              target_cols,
              layers=("auto", i, 10, 1),
              activ_func="tanh",
              epochs=1000,
              l_rate=0.3
          )
```

```
[33]: for i in range(1,11):
      print(f"Model Trained with Hidden Layer Configuration ({i}, 10)",
            end=f"\n{'-'*100}\n")

      print("Final Training Results")
      print(layer_tests_3[f"Test-{i}"]["training_results"].iloc[-1, :4],
            end=f"\n{'-'*100}\n")

      print("Final Validation Results")
      print(layer_tests_3[f"Test-{i}"]["training_results"].iloc[-1, 8:12],
            end=f"\n{'-'*100}\n")

      print("Test Set Results")
      print(layer_tests_3[f"Test-{i}"]["error_metrics"].iloc[0],
            end=f"\n{'-'*100}\n\n\n")

      ax = layer_tests_3[f"Test-{i}"]["training_results"].plot(
          y=["rmse", "val_rmse"], title=f"Model Trained with Hidden Layer
            Configuration ({i}, 10)",
      )
      ax.set_xlabel("Epochs")
      ax.set_ylabel("Root Mean Squared Error (RMSE)")
```

Model Trained with Hidden Layer Configuration (1, 10)

Final Training Results

mae	40.675349
mse	3289.605540
r_sqr	0.000141
rmse	57.355083
Name: 999, dtype: float64	

Final Validation Results

```
val_mae      40.452444
val_mse      3191.763601
val_r_sqr    -0.000552
val_rmse     56.495695
Name: 999, dtype: float64
```


Test Set Results

```
mse          2442.108068
rmse         49.417690
mae          36.684681
r_sqr        -0.005315
st_mse       0.024735
st_rmse      0.157273
st_mae       0.094491
st_r_sqr     -0.084921
Name: 0, dtype: float64
```

=====

Model Trained with Hidden Layer Configuration (2, 10)

Final Training Results

```
mae          9.205754
mse          279.591917
r_sqr        0.907125
rmse         16.721002
Name: 999, dtype: float64
```


Final Validation Results

```
val_mae      10.477098
val_mse      355.636679
val_r_sqr    0.903670
val_rmse     18.858332
Name: 999, dtype: float64
```


Test Set Results

```
mse          948.689127
rmse         30.800798
mae          19.482716
r_sqr        0.655228
st_mse       0.002617
st_rmse      0.051161
```

```
st_mae      0.023390
st_r_sqr     0.886004
Name: 0, dtype: float64
```

```
=====
=====
```

Model Trained with Hidden Layer Configuration (3, 10)

```
-----
-----
```

Final Training Results

```
mae      29.776235
mse     1000.587635
r_sqr     0.651198
rmse     31.632067
Name: 999, dtype: float64
```

```
-----
-----
```

Final Validation Results

```
val_mae     30.532874
val_mse     1100.822412
val_r_sqr    0.716457
val_rmse     33.178644
Name: 999, dtype: float64
```

```
-----
-----
```

Test Set Results

```
mse      1227.786313
rmse      35.039782
mae       30.576602
r_sqr     0.589698
st_mse     0.003323
st_rmse    0.057648
st_mae     0.050783
st_r_sqr    0.658818
Name: 0, dtype: float64
```

```
=====
=====
```

Model Trained with Hidden Layer Configuration (4, 10)

```
-----
-----
```

Final Training Results

```
mae      9.533718
mse     299.505396
```

```
r_sqr      0.901843
rmse      17.306224
Name: 999, dtype: float64
```

Final Validation Results

```
val_mae      8.655035
val_mse     241.227884
val_r_sqr    0.908075
val_rmse     15.531513
Name: 999, dtype: float64
```

Test Set Results

```
mse      545.763014
rmse     23.361571
mae      12.153721
r_sqr    0.852945
st_mse    0.001061
st_rmse   0.032575
st_mae    0.018374
st_r_sqr  0.911921
Name: 0, dtype: float64
```

Model Trained with Hidden Layer Configuration (5, 10)

Final Training Results

```
mae      8.727363
mse     216.310404
r_sqr    0.928333
rmse     14.707495
Name: 999, dtype: float64
```

Final Validation Results

```
val_mae      9.774685
val_mse     322.415189
val_r_sqr    0.907832
val_rmse     17.955924
Name: 999, dtype: float64
```

Test Set Results


```

mse          487.227778
rmse         22.073237
mae          14.595240
r_sqr        0.834330
st_mse       0.002324
st_rmse      0.048203
st_mae       0.023174
st_r_sqr     0.878759
Name: 0, dtype: float64
=====
=====

```

Model Trained with Hidden Layer Configuration (6, 10)

```

-----
Final Training Results
mae          6.903358
mse         159.462782
r_sqr        0.952458
rmse         12.627857
Name: 999, dtype: float64
-----

```

```

-----
Final Validation Results
val_mae       6.846378
val_mse      169.339608
val_r_sqr     0.950305
val_rmse     13.013055
Name: 999, dtype: float64
-----

```

```

-----
Test Set Results
mse          2610.194233
rmse         51.090060
mae          35.145469
r_sqr       -0.311430
st_mse       0.002085
st_rmse      0.045664
st_mae       0.022646
st_r_sqr     0.922044
Name: 0, dtype: float64
=====
=====

```

Model Trained with Hidden Layer Configuration (7, 10)

Final Training Results

```
mae      7.294836
mse     166.755323
r_sqr    0.948493
rmse    12.913378
Name: 999, dtype: float64
```

Final Validation Results

```
val_mae    7.466487
val_mse   188.788827
val_r_sqr   0.936254
val_rmse   13.740045
Name: 999, dtype: float64
```

Test Set Results

```
mse      246.147520
rmse     15.689089
mae       7.769502
r_sqr     0.912560
st_mse     0.001219
st_rmse    0.034917
st_mae     0.019874
st_r_sqr   0.866350
Name: 0, dtype: float64
```

Model Trained with Hidden Layer Configuration (8, 10)

Final Training Results

```
mae      6.053548
mse     124.757734
r_sqr    0.961100
rmse    11.169500
Name: 999, dtype: float64
```

Final Validation Results

```
val_mae    5.847862
val_mse   102.754911
```

```
val_r_sqr      0.961203
val_rmse       10.136810
Name: 999, dtype: float64
```

Test Set Results

```
mse           2323.555325
rmse          48.203271
mae           34.414942
r_sqr         0.279108
st_mse        0.002943
st_rmse       0.054247
st_mae        0.036393
st_r_sqr      0.857782
Name: 0, dtype: float64
```

Model Trained with Hidden Layer Configuration (9, 10)

Final Training Results

```
mae           8.399457
mse          218.070847
r_sqr        0.933390
rmse         14.767222
Name: 999, dtype: float64
```

Final Validation Results

```
val_mae       7.389402
val_mse      157.882046
val_r_sqr     0.941587
val_rmse     12.565112
Name: 999, dtype: float64
```

Test Set Results

```
mse           706.22259
rmse          26.574843
mae           16.900812
r_sqr         0.761802
st_mse        0.001696
st_rmse       0.041186
st_mae        0.021709
st_r_sqr      0.910816
```

Name: 0, dtype: float64

=====

Model Trained with Hidden Layer Configuration (10, 10)

Final Training Results

mae 8.578838

mse 254.414124

r_sqr 0.913281

rmse 15.950364

Name: 999, dtype: float64

Final Validation Results

val_mae 8.812246

val_mse 340.880065

val_r_sqr 0.923956

val_rmse 18.462938

Name: 999, dtype: float64

Test Set Results

mse 225.380567

rmse 15.012680

mae 9.438959

r_sqr 0.896025

st_mse 0.005652

st_rmse 0.075180

st_mae 0.042396

st_r_sqr 0.721736

Name: 0, dtype: float64

=====

=====

